

# International Workshop on Machine Learning for Space Weather: Fundamentals, Tools and Future Prospects

**7-11 November 2022**  
**This is a hybrid meeting**  
**Buenos Aires, Argentina**



Further information:

<https://indico.ictp.it/event/9840/>

smr3750@ictp.it

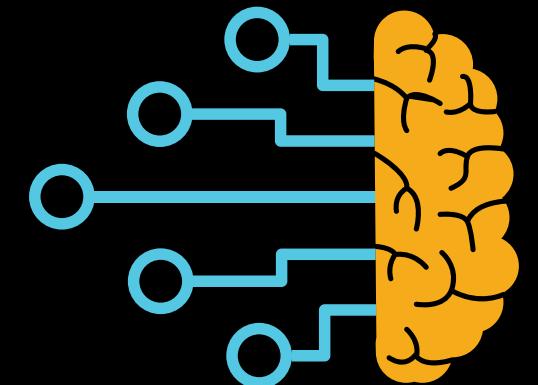
+39-040-2240284

Elizabeth Brancaccio

**Dra María Graciela Molina**  
FACET-UNT / CONICET  
Tucumán Space Weather Center - TSWC

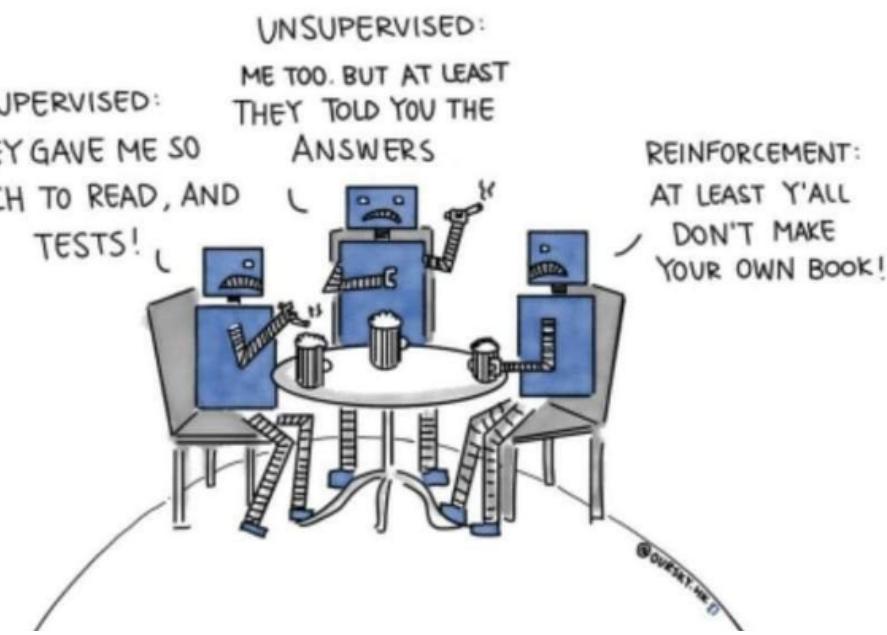
<https://spaceweather.facet.unt.edu.ar/>  
IG -> @spaceweatherargentina

gmolina@herrera.unt.edu.ar



# Algorithms

# MACHINE LEARNING



## Supervised



### Regression

- Linear
- Polynomial

## Unsupervised



### Clustering

- K-means
- PCA
- SVD

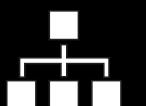
Continuous  
data



### Decision Tree



### Random Forest



### Classification

- KNN
- Logistic Regression
- Naive-Bayes
- SVM

Association Rule  
Learning

Categorical  
data

Hidden Markov  
Model

## Artificial Neural Networks



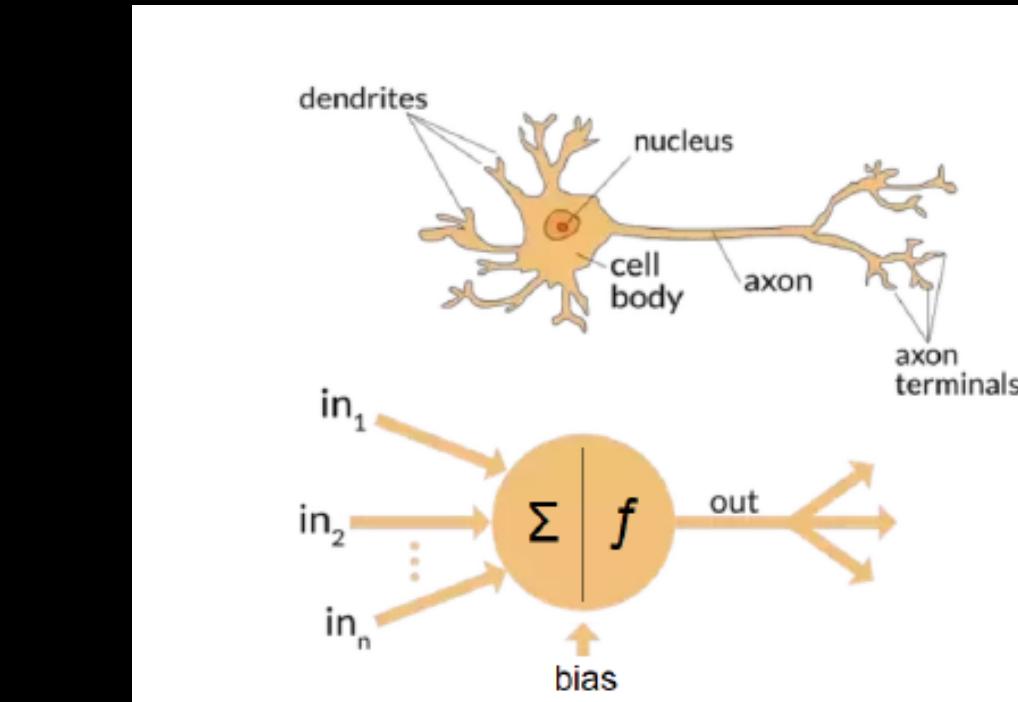
TSWC, 2022

# Artificial Neural Networks

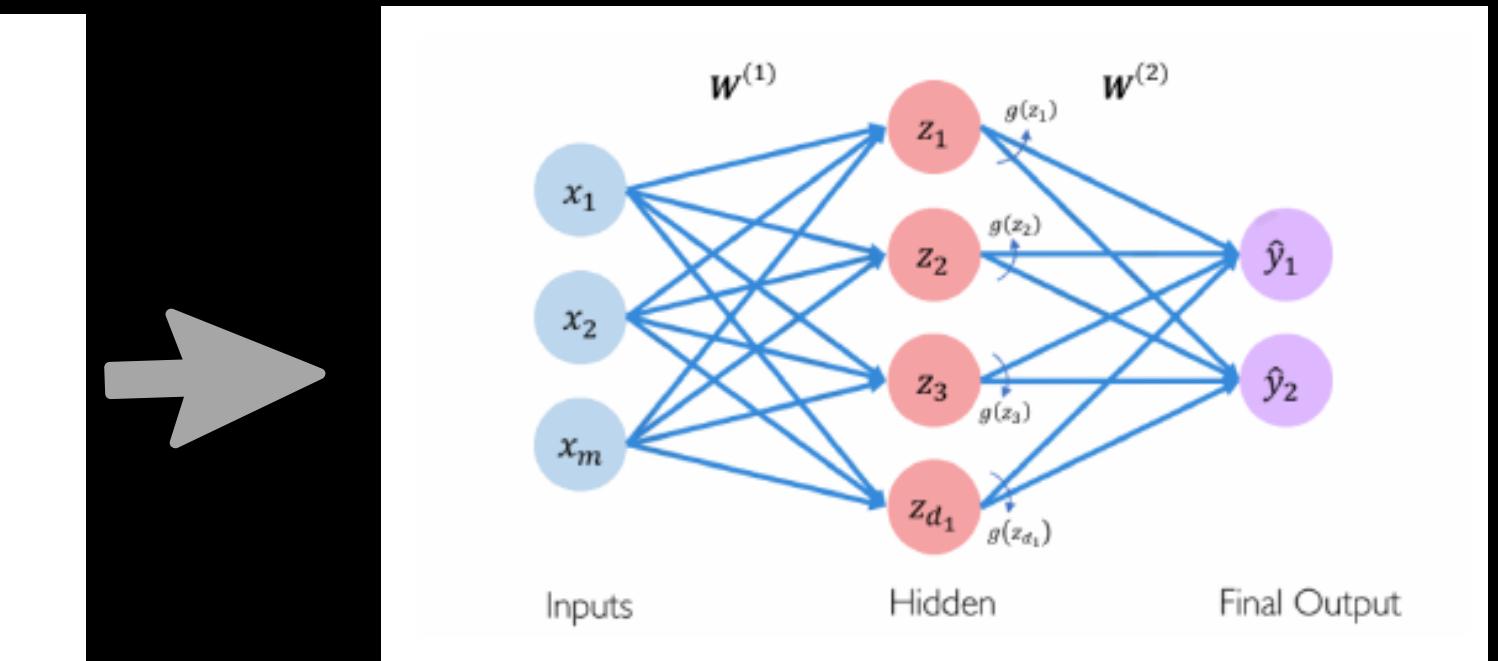
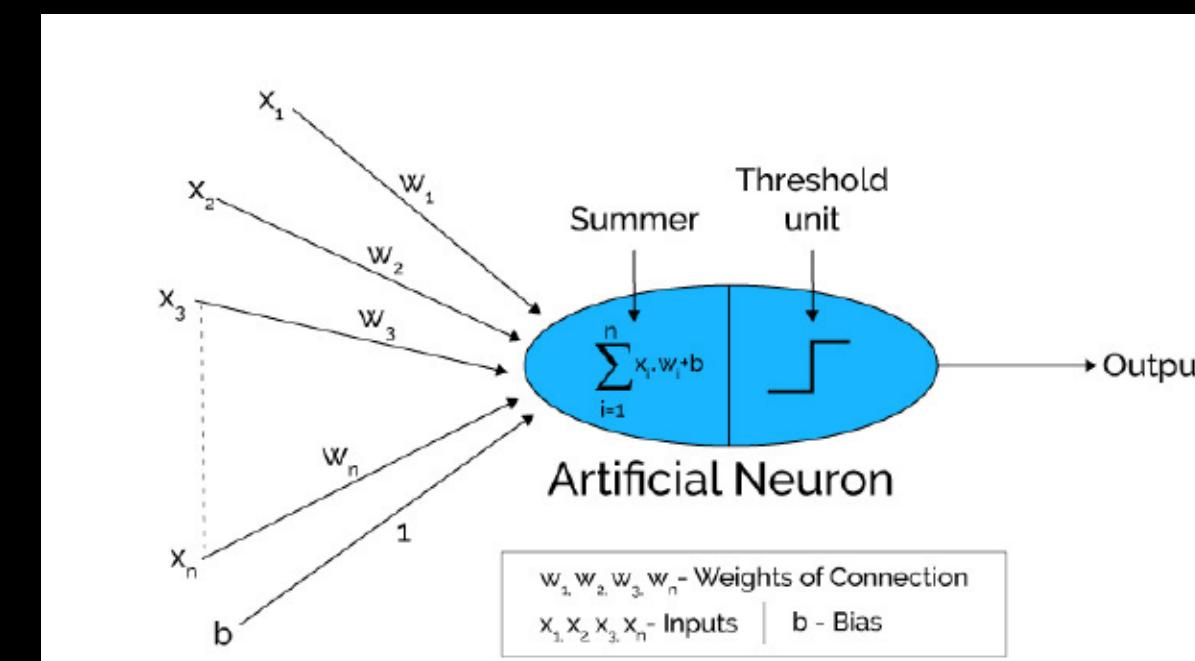


TSWC, 2022

ANN



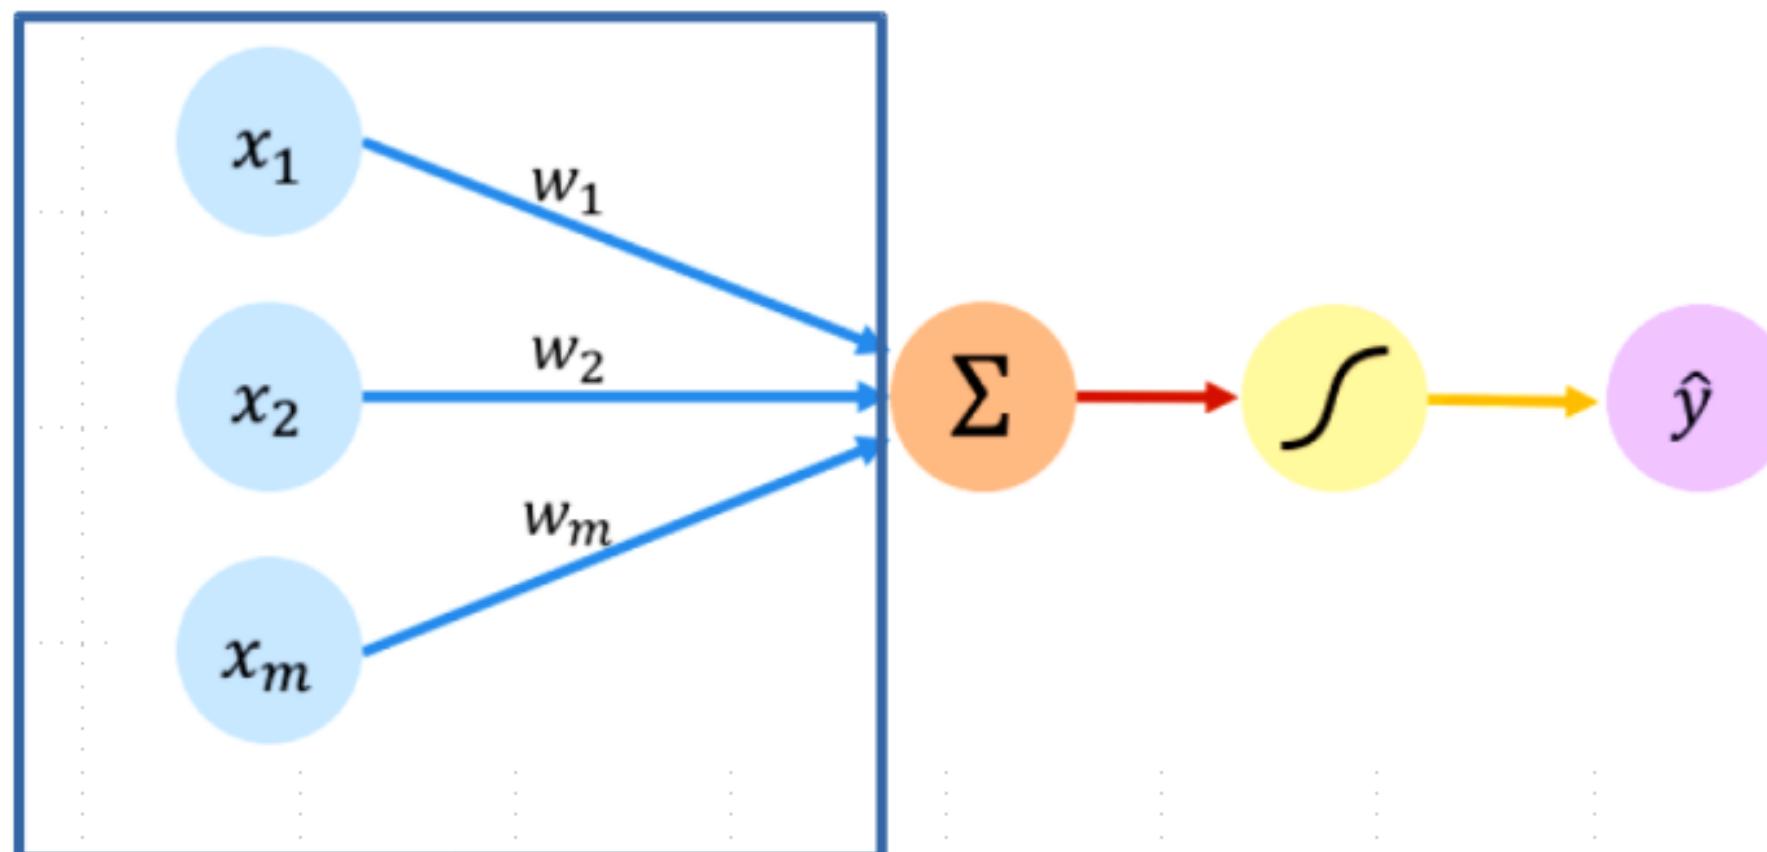
- Data-driven modeling
- Inspired in brain neural networks
- Solving wide number of complex problems: facial recognition, handwrite recognition, weather forecasting, etc.
- Black-box modelling (?) based in the composition of interconnected neurons (neural network)



# Perceptron

- Unidad fundamental para construir redes neuronales

Input → Weights → Sum → Non linearity → Output

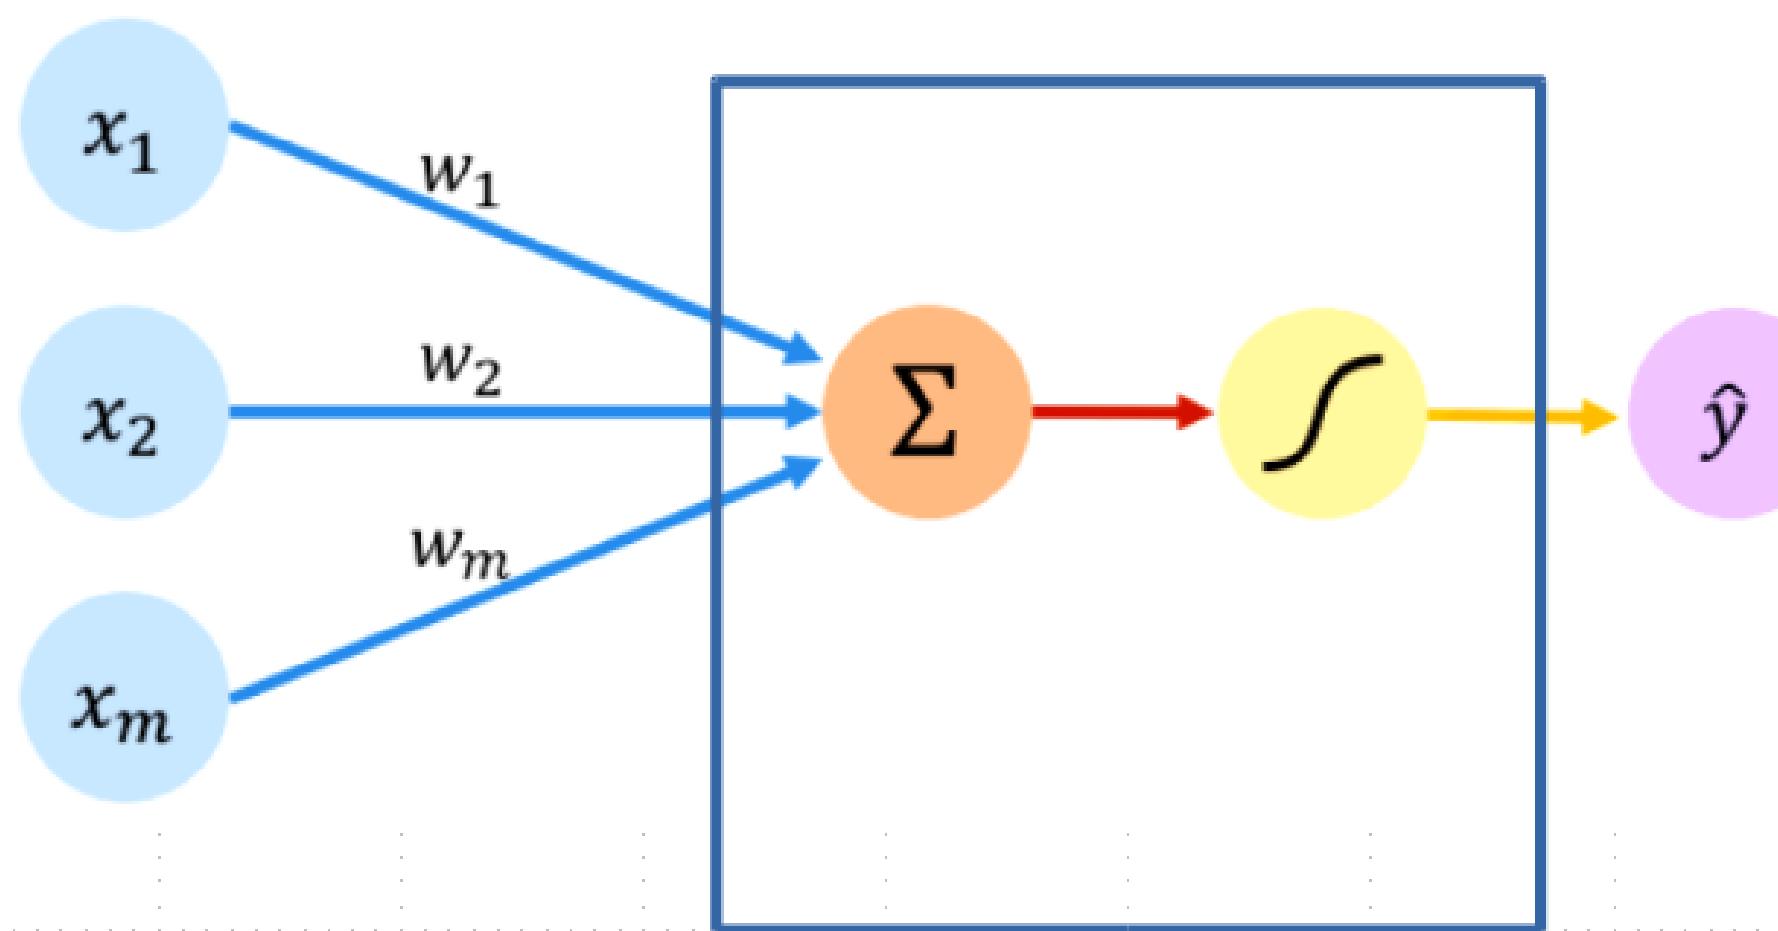


- Cada una de las entradas corresponde a una característica (feature)
- A cada una de las  $x_i$  se les aplica un peso ( $w$ ) ( ponderación de los valores de la entrada)
- Los pesos son las variables que se ajustan durante el entrenamiento



# Perceptron

Input → Weights → Sum → Non linearity → Output



Linear combination  
of inputs

$$\hat{y} = g \left( w_0 + \sum_{i=1}^m x_i w_i \right)$$

Output

Non-linear activation function

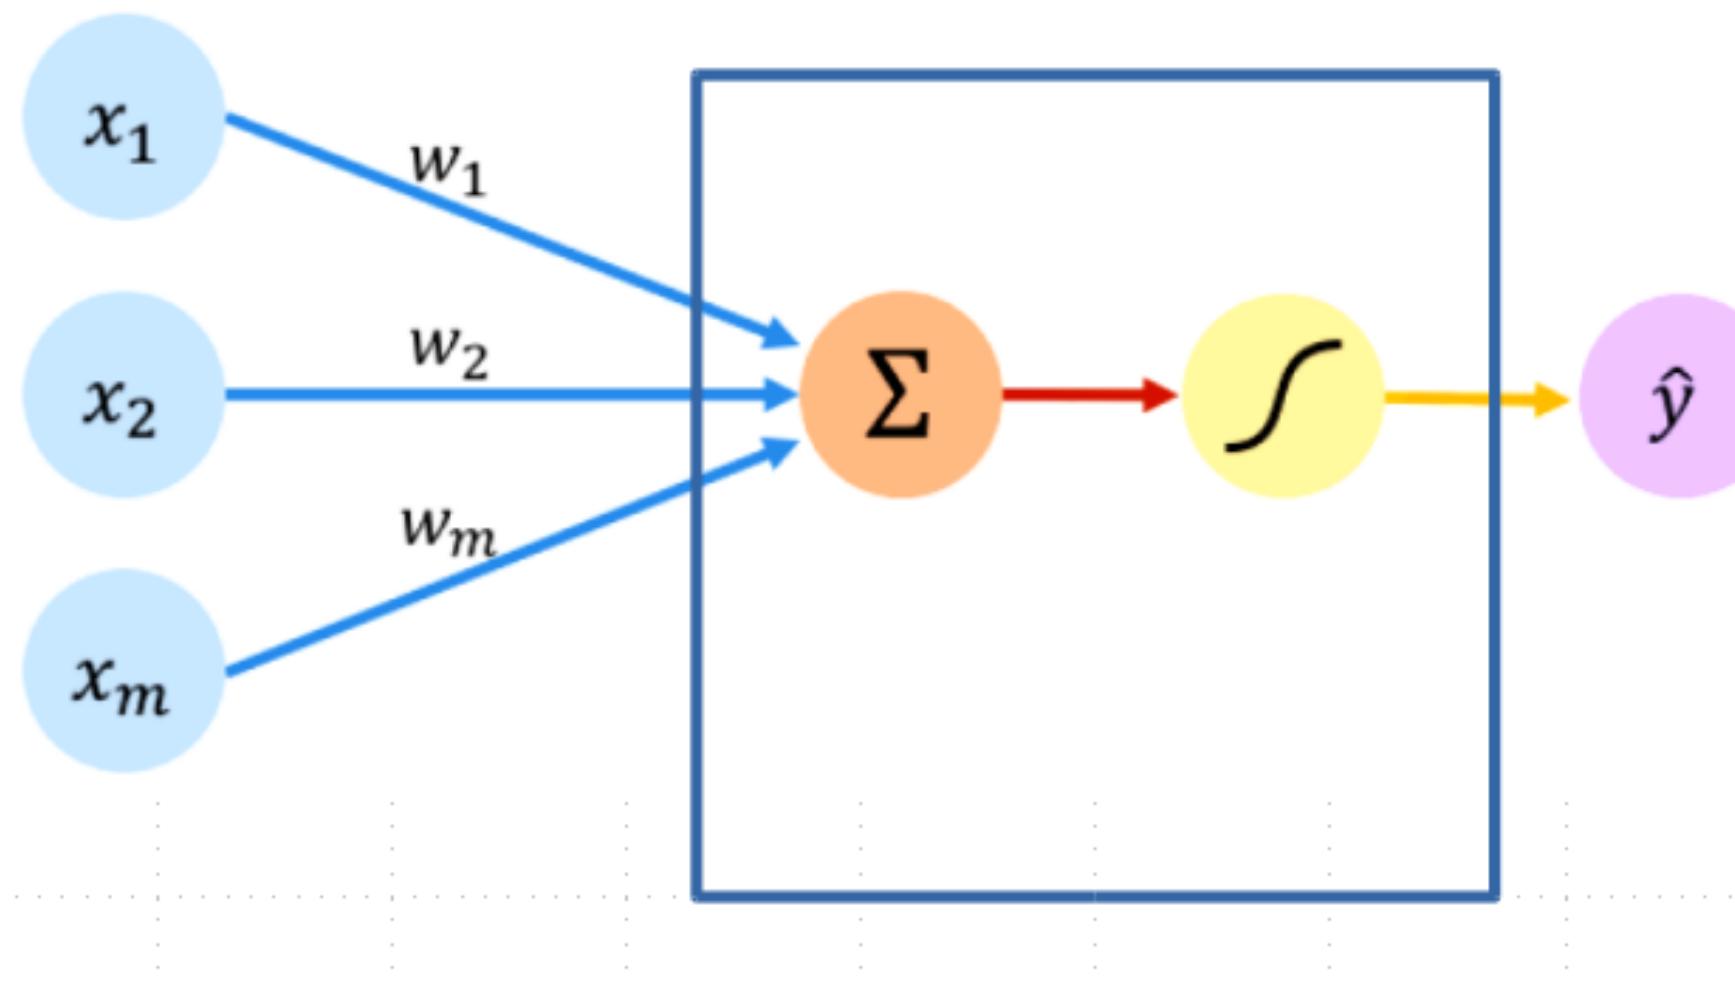
Bias

- Capa oculta (hidden layer).
- El bias es un término de ajuste
- La función de activación se aplica a la combinación lineal de las entradas ponderadas teniendo en cuenta el bias



# Perceptron

Input → Weights → Sum → Non linearity → Output



$$\hat{y} = g \left( w_0 + \sum_{i=1}^m x_i w_i \right)$$

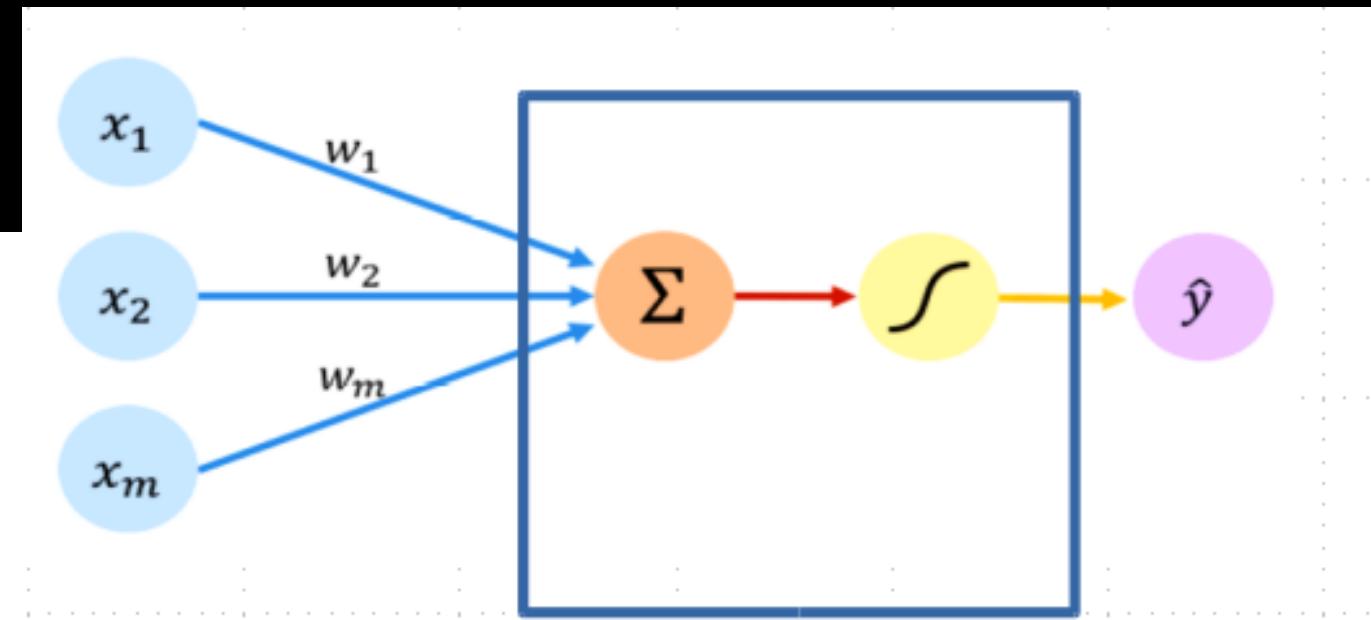
$$\hat{y} = g ( w_0 + \mathbf{X}^T \mathbf{W} )$$

where:  $\mathbf{X} = \begin{bmatrix} x_1 \\ \vdots \\ x_m \end{bmatrix}$  and  $\mathbf{W} = \begin{bmatrix} w_1 \\ \vdots \\ w_m \end{bmatrix}$



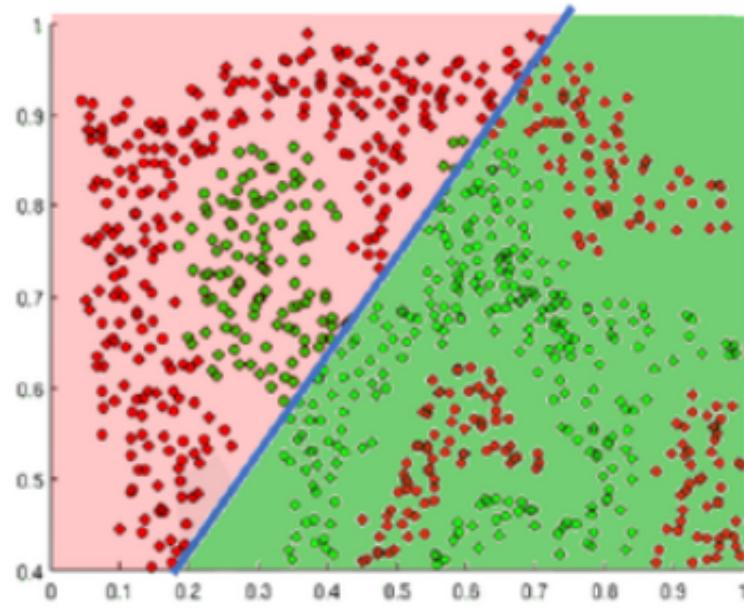
# Activation function

- Agrega no-linealidad al modelo

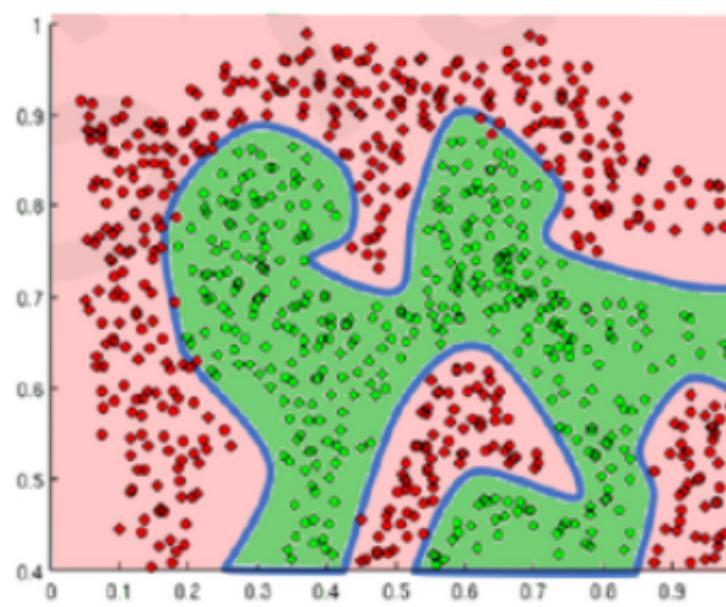


$$\hat{y} = g(w_0 + X^T W)$$

- Una de las más usadas es la función sigmoide.
- Produce resultados entre 0 y 1

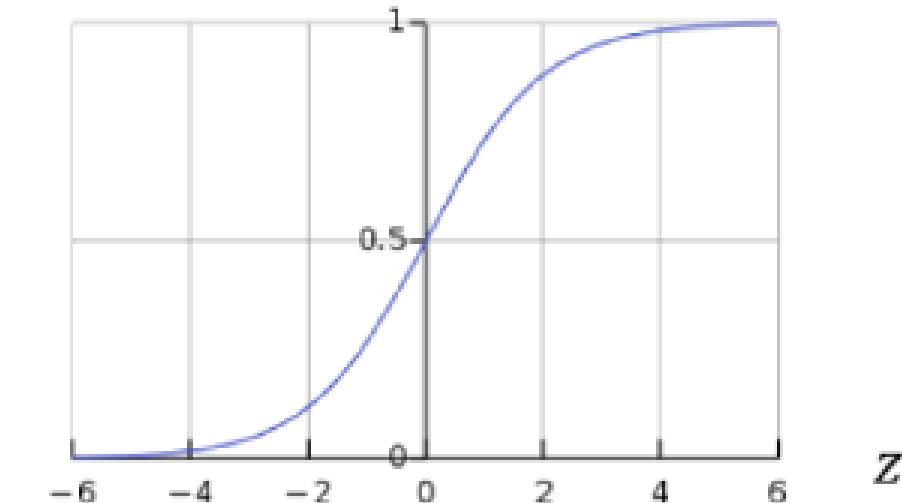


Linear activation functions produce linear decisions no matter the network size



Non-linearities allow us to approximate arbitrarily complex functions

$$g(z) = \sigma(z) = \frac{1}{1 + e^{-z}}$$

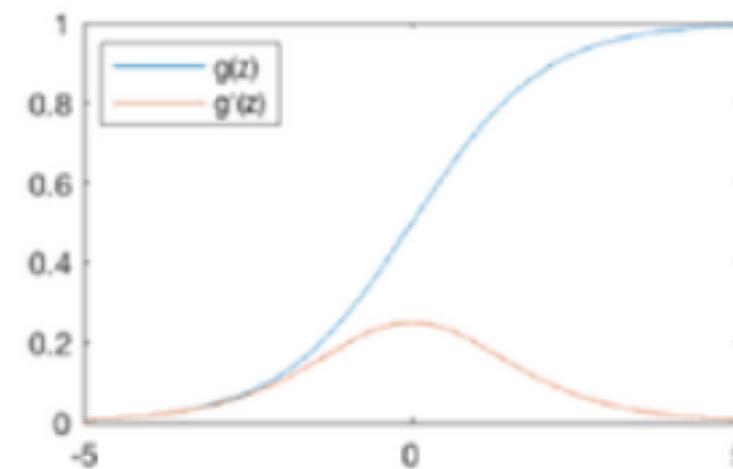




# Activation function

$$\hat{y} = g(w_0 + X^T W)$$

Sigmoid Function

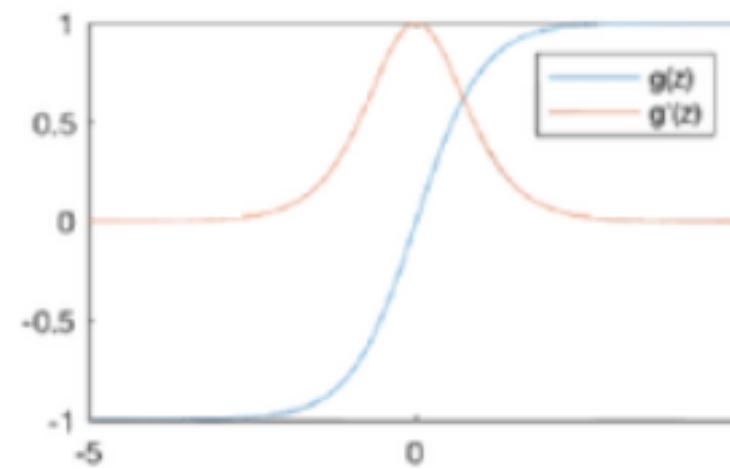


$$g(z) = \frac{1}{1 + e^{-z}}$$

$$g'(z) = g(z)(1 - g(z))$$

`tf.math.sigmoid(z)`

Hyperbolic Tangent

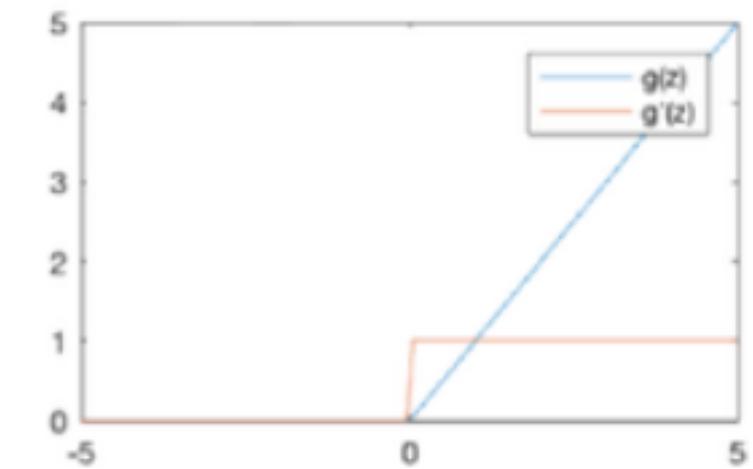


$$g(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

$$g'(z) = 1 - g(z)^2$$

`tf.math.tanh(z)`

Rectified Linear Unit (ReLU)



$$g(z) = \max(0, z)$$

$$g'(z) = \begin{cases} 1, & z > 0 \\ 0, & \text{otherwise} \end{cases}$$

`tf.nn.relu(z)`

# Activation function

- Adds no-linearity to the model

$$\hat{y} = g(w_0 + X^T W)$$

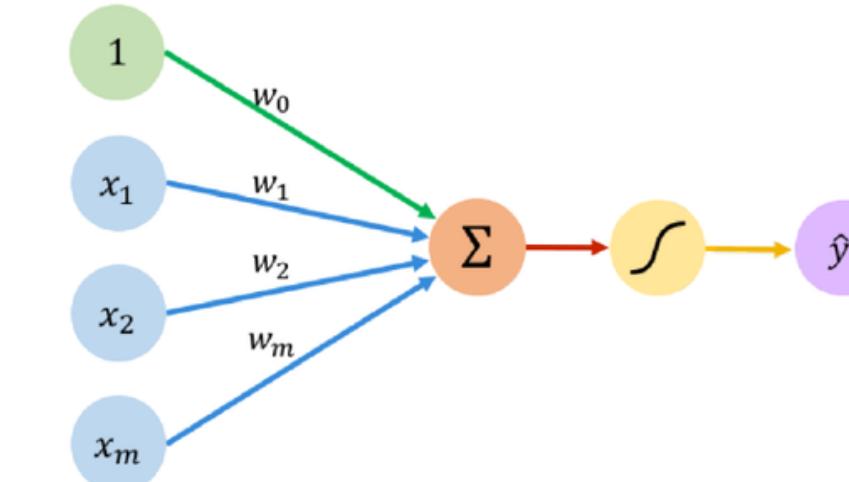
Name	Plot	Equation	Derivative
Identity		$f(x) = x$	$f'(x) = 1$
Binary step		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x \neq 0 \\ ? & \text{for } x = 0 \end{cases}$
Logistic (a.k.a Soft step)		$f(x) = \frac{1}{1 + e^{-x}}$	$f'(x) = f(x)(1 - f(x))$
Tanh		$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$	$f'(x) = 1 - f(x)^2$
Arctan		$f(x) = \tan^{-1}(x)$	$f'(x) = \frac{1}{x^2 + 1}$
Rectified Linear Unit (ReLU)		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Parameteric Rectified Linear Unit (PReLU) [2]		$f(x) = \begin{cases} \alpha x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Exponential Linear Unit (ELU) [3]		$f(x) = \begin{cases} \alpha(e^x - 1) & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} f(x) + \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
SoftPlus		$f(x) = \log_e(1 + e^x)$	$f'(x) = \frac{1}{1 + e^{-x}}$



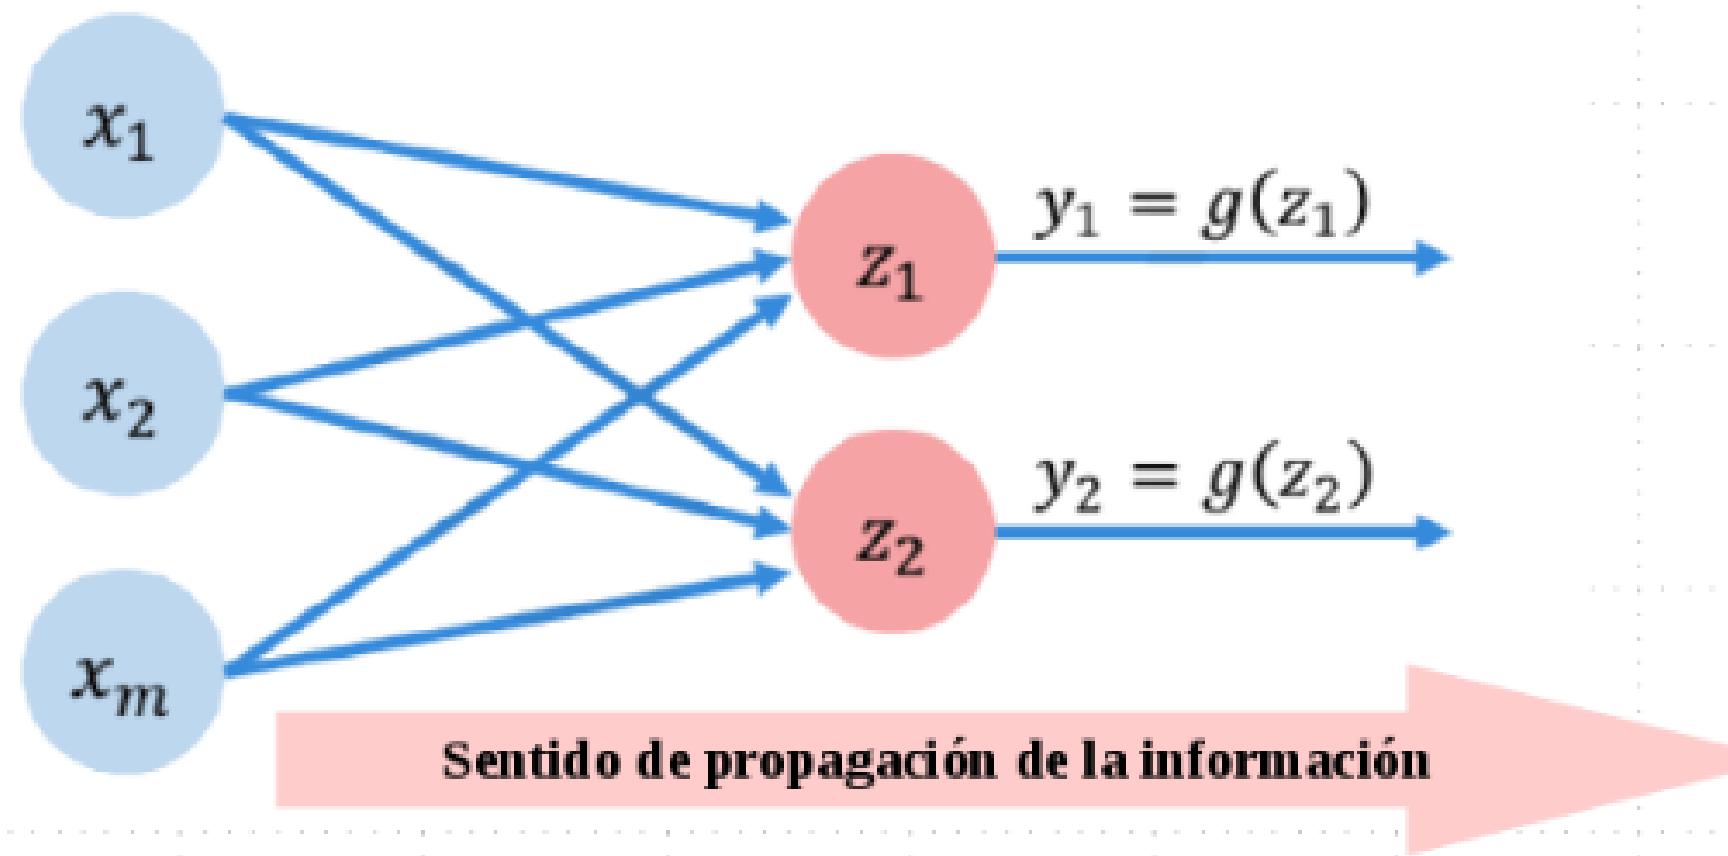
# Feed Forward (FF)

- Perceptrón simplificado (n entradas → 1 salida)

$$z = w_0 + \sum_{j=1}^m x_j w_j$$



- Armemos una red neuronal compuesta por múltiples perceptrones (n entradas → m salida)

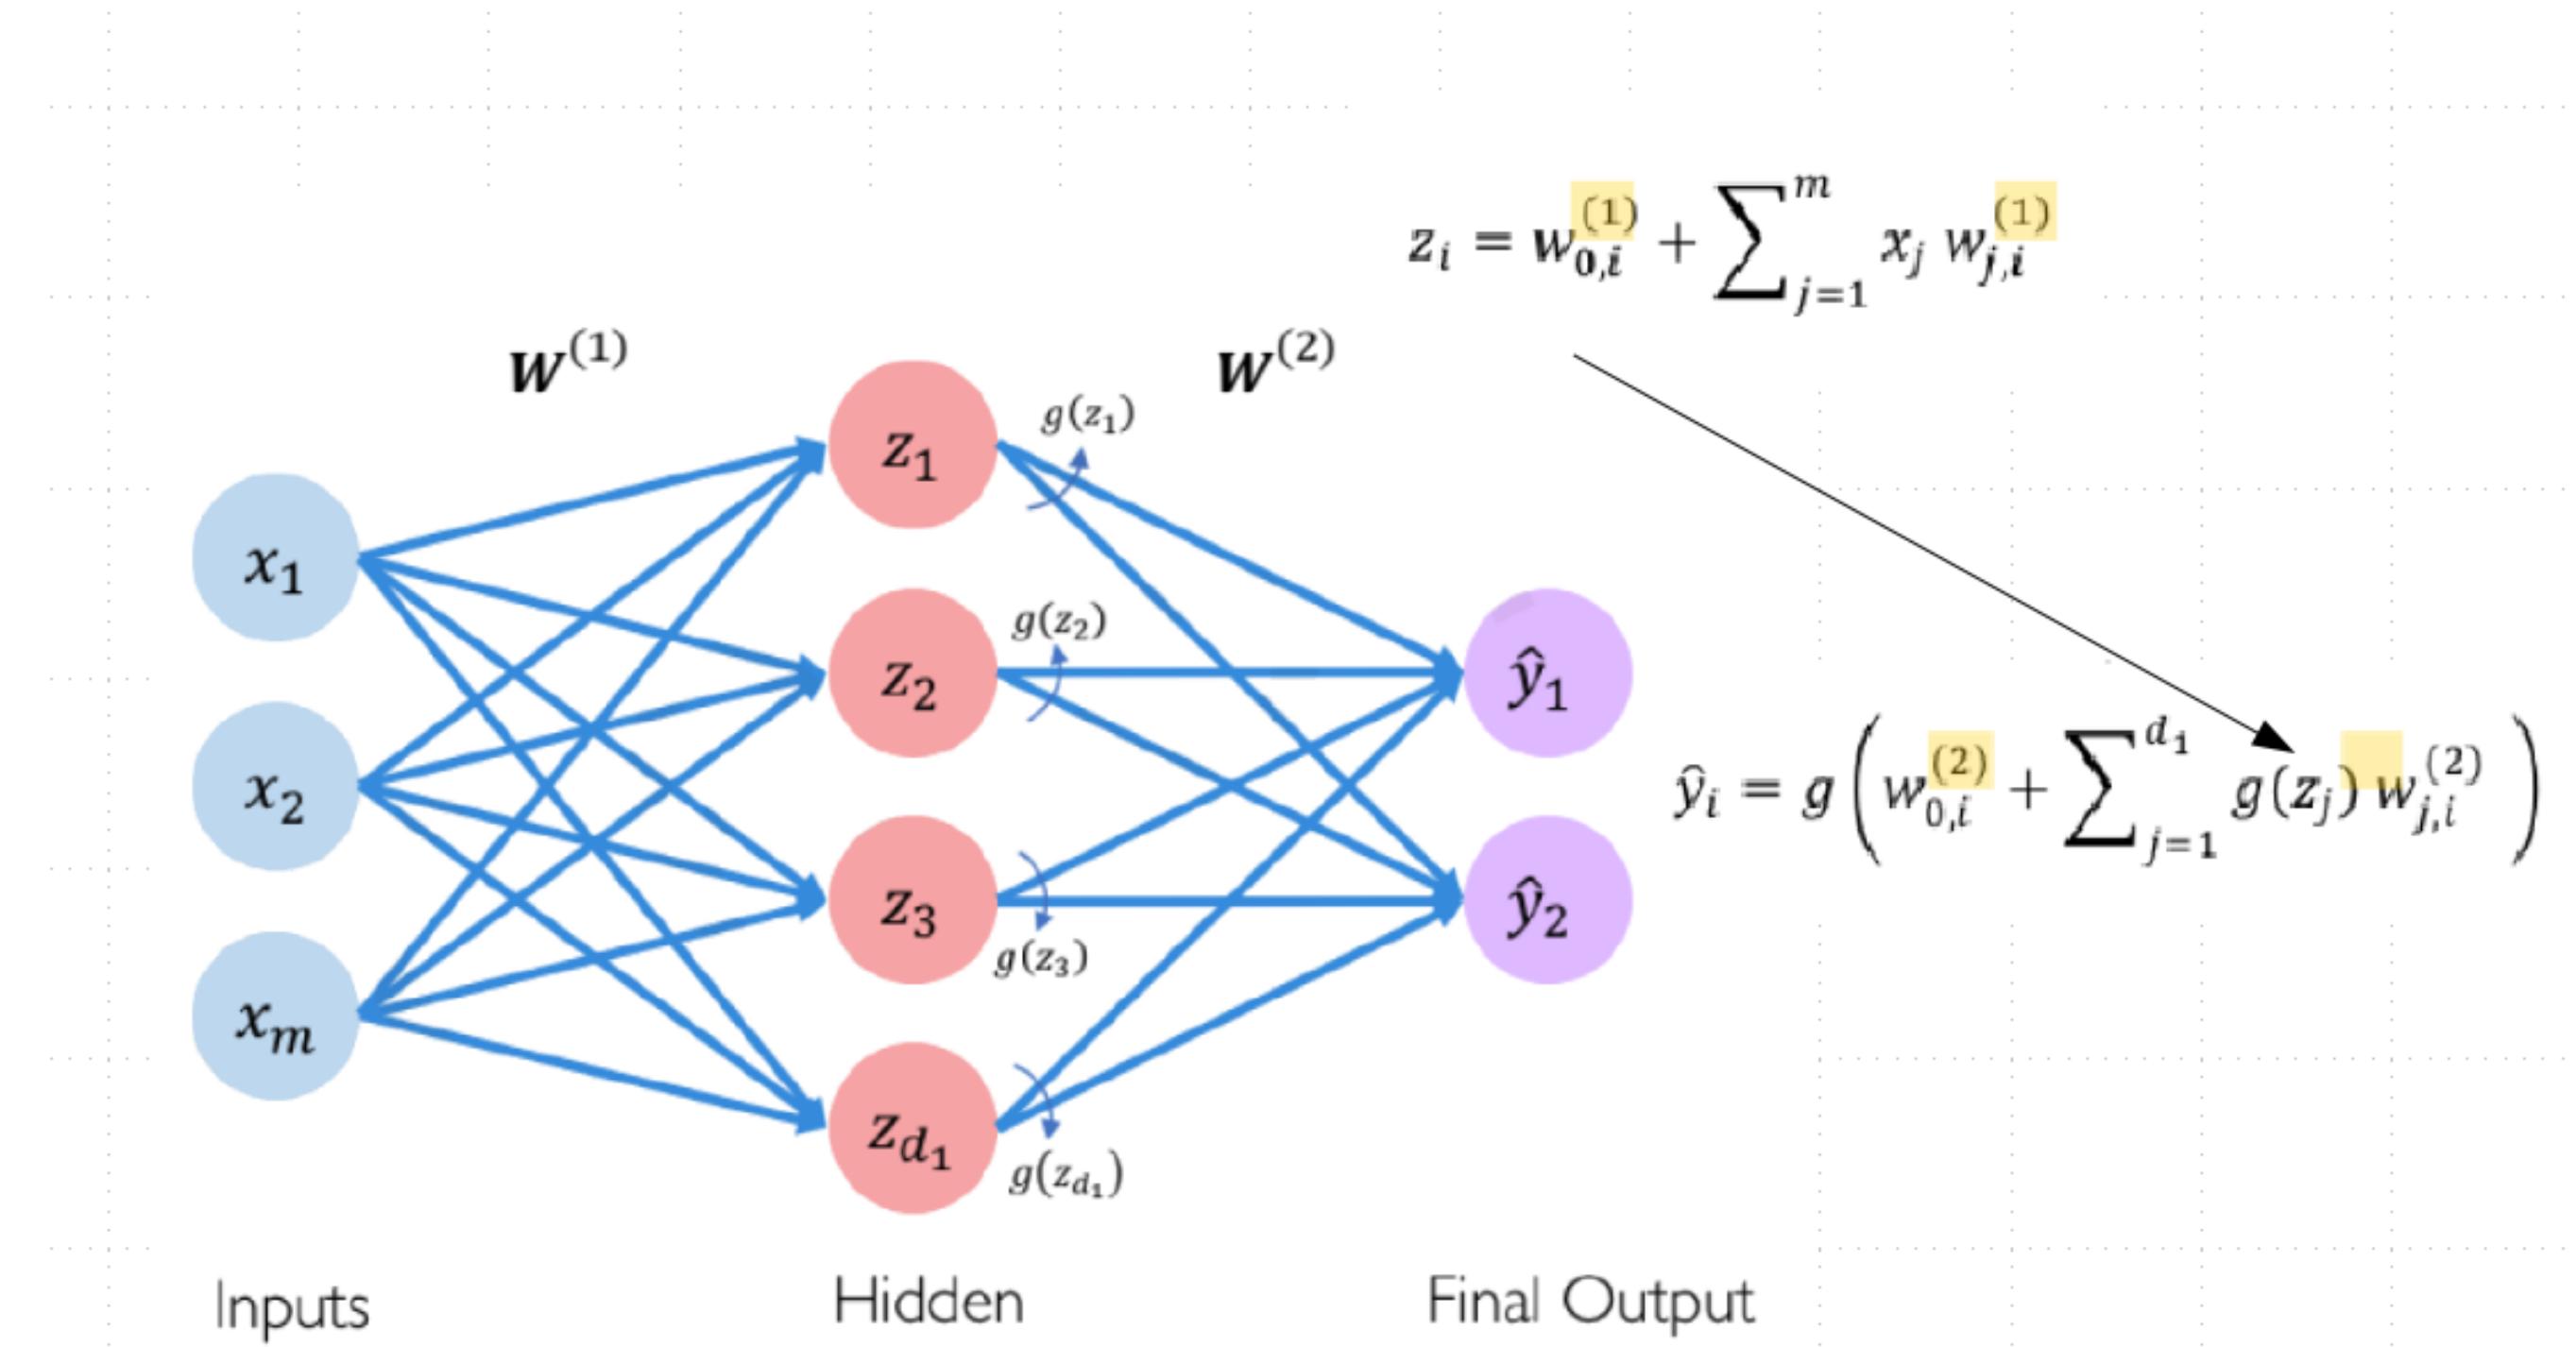


$$z_i = w_{0,i} + \sum_{j=1}^m x_j w_{j,i}$$

Observación: todas las entradas están conectadas a todas las salidas → dense layers

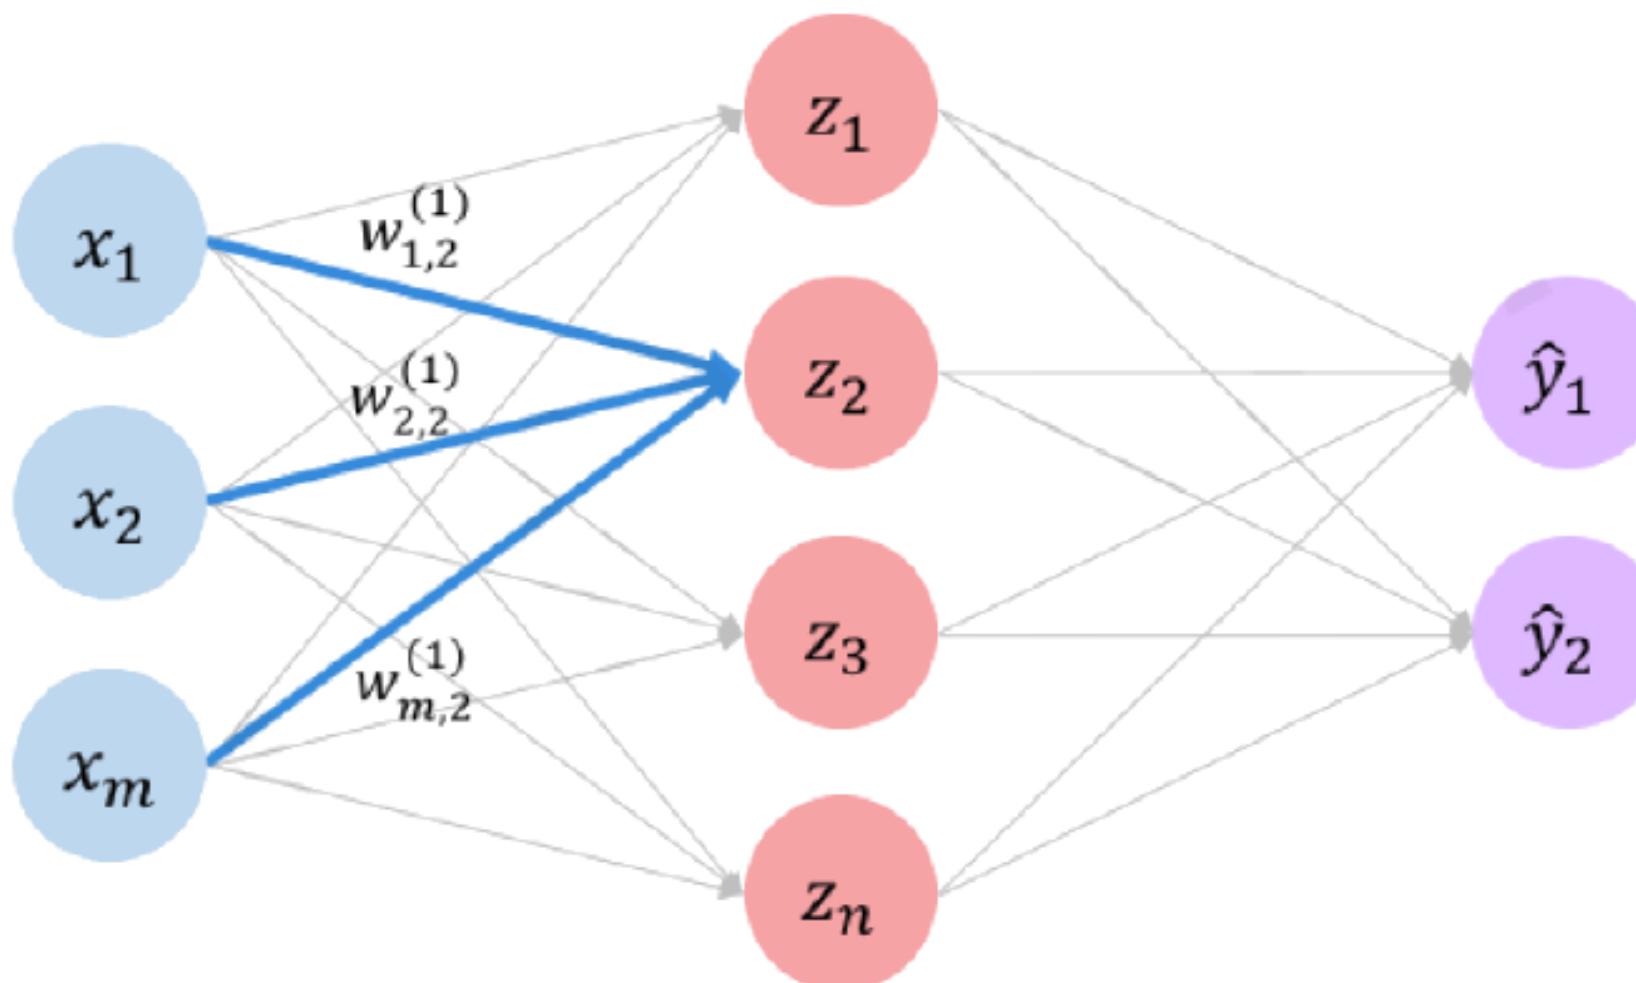


# 1-Layer Neural Network





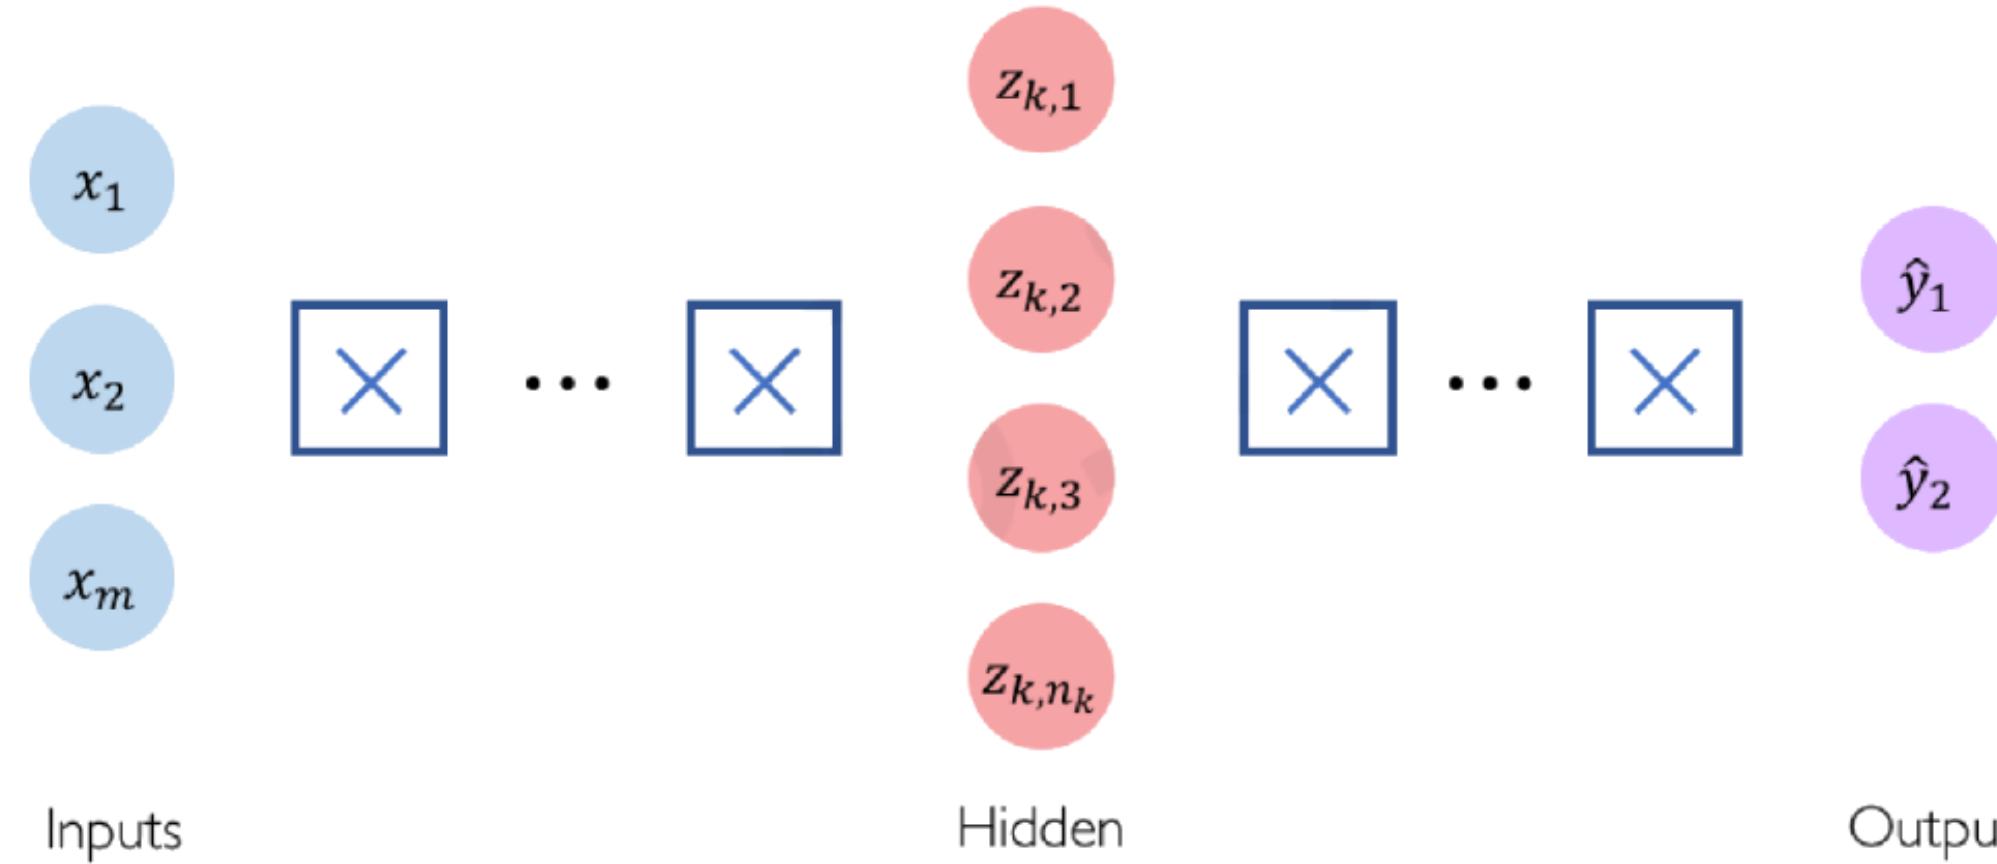
# 1-Layer Neural Network



$$\begin{aligned} z_2 &= w_{0,2}^{(1)} + \sum_{j=1}^m x_j w_{j,2}^{(1)} \\ &= w_{0,2}^{(1)} + x_1 w_{1,2}^{(1)} + x_2 w_{2,2}^{(1)} + x_m w_{m,2}^{(1)} \end{aligned}$$

# M-layers Neural Network (multi-layer)

- Deep Neural Network

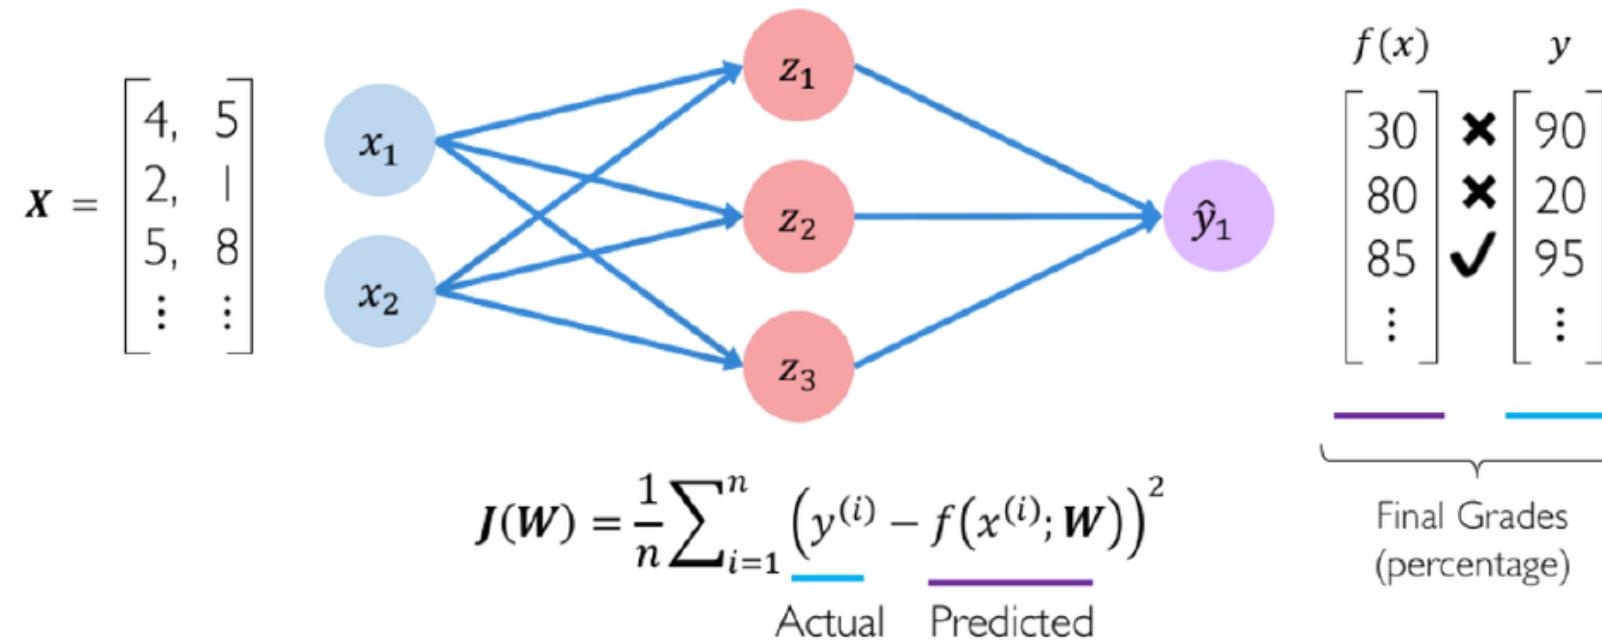


$$z_{k,i} = w_{0,i}^{(k)} + \sum_{j=1}^{n_{k-1}} g(z_{k-1,j}) w_{j,i}^{(k)}$$



# Loss Function

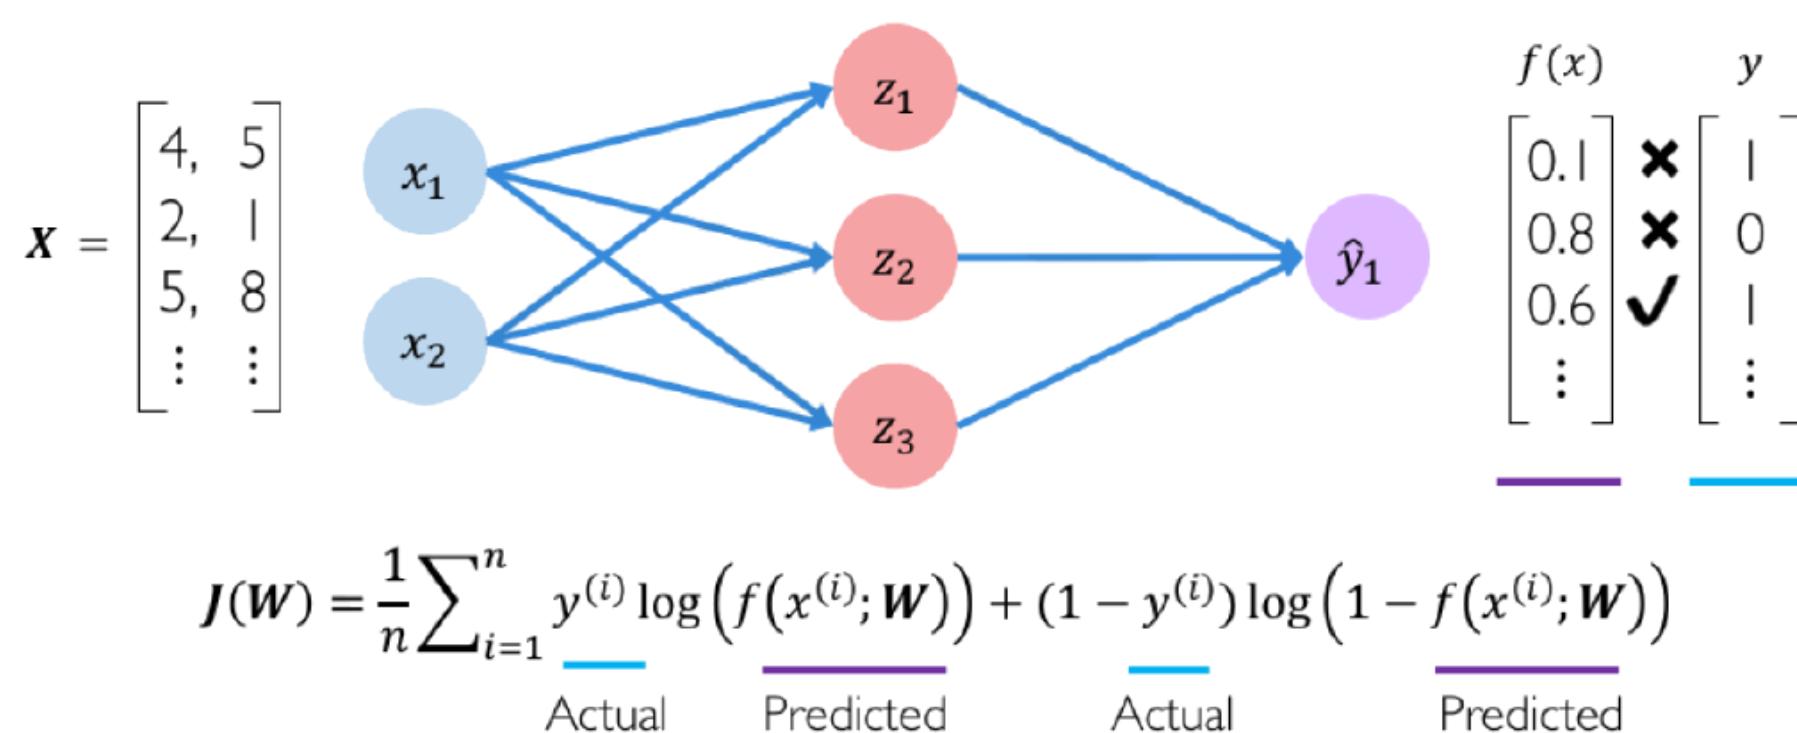
Mean squared error loss can be used with regression models that output continuous real numbers



$f(x)$	$y$
30	✗ 90
80	✗ 20
85	✓ 95
⋮	⋮

Final Grades (percentage)

Cross entropy loss can be used with models that output a probability between 0 and 1



symbol	name	equation
$\mathcal{L}_1$	$L_1$ loss	$\ y - o\ _1$
$\mathcal{L}_2$	$L_2$ loss	$\ y - o\ _2^2$
$\mathcal{L}_1 \circ \sigma$	expectation loss	$\ y - \sigma(o)\ _1$
$\mathcal{L}_2 \circ \sigma$	regularised expectation loss <sup>1</sup>	$\ y - \sigma(o)\ _2^2$
$\mathcal{L}_\infty \circ \sigma$	Chebyshev loss	$\max_j  \sigma(o)^{(j)} - y^{(j)} $
hinge	hinge [13] (margin) loss	$\sum_j \max(0, \frac{1}{2} - \hat{y}^{(j)} o^{(j)})$
hinge <sup>2</sup>	squared hinge (margin) loss	$\sum_j \max(0, \frac{1}{2} - \hat{y}^{(j)} o^{(j)})^2$
hinge <sup>3</sup>	cubed hinge (margin) loss	$\sum_j \max(0, \frac{1}{2} - \hat{y}^{(j)} o^{(j)})^3$
log	log (cross entropy) loss	$-\sum_j y^{(j)} \log \sigma(o^{(j)})$
log <sup>2</sup>	squared log loss	$-\sum_j [y^{(j)} \log \sigma(o^{(j)})]^2$
tan	Tanimoto loss	$-\sum_j \sigma(o)^{(j)} y^{(j)}$
D <sub>CS</sub>	Cauchy-Schwarz Divergence [3]	$\frac{\ \sigma(o)\ _2^2 + \ y\ _2^2 - \sum_j \sigma(o)^{(j)} y^{(j)}}{\ \sigma(o)\ _2 \ y\ _2}$

- <https://arxiv.org/pdf/1702.05659.pdf>





# Training (loss optimization)

- Encontrar los valores de los pesos de la red neuronal tal que la pérdida o costo sea mínimo

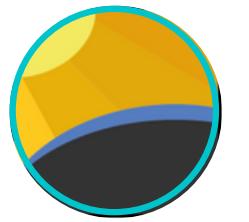
—

$$\mathbf{W}^* = \operatorname{argmin}_{\mathbf{W}} \frac{1}{n} \sum_{i=1}^n \mathcal{L}(f(x^{(i)}; \mathbf{W}), y^{(i)})$$

$$\mathbf{W}^* = \operatorname{argmin}_{\mathbf{W}} J(\mathbf{W})$$

- Es un problema de optimización
- Podemos usar algún método numérico de optimización

MÉTODO DEL GRADIENTE DESCENDIENTE



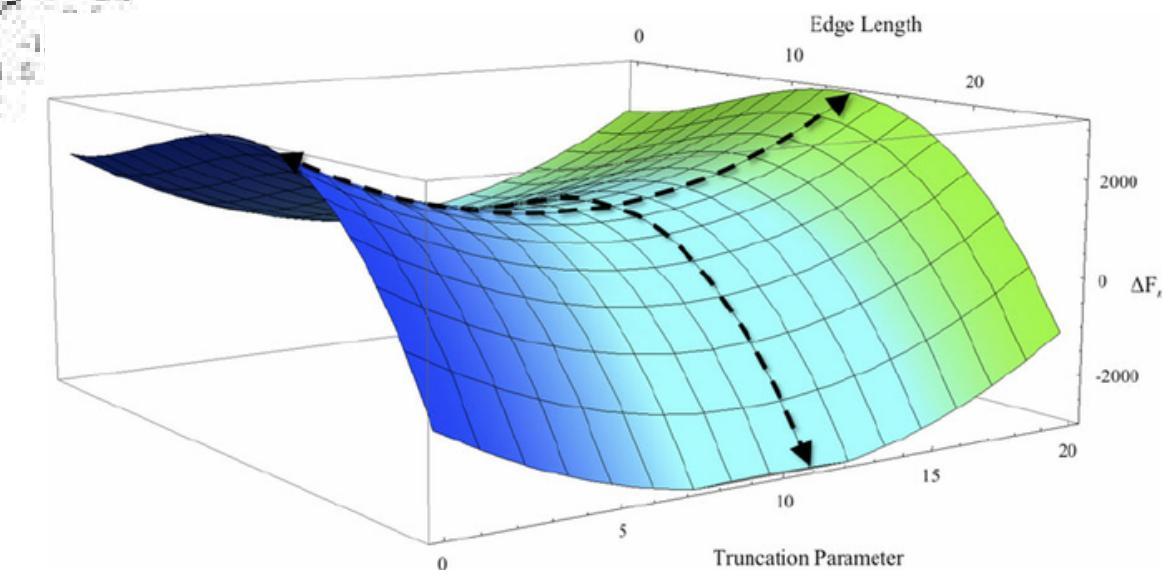
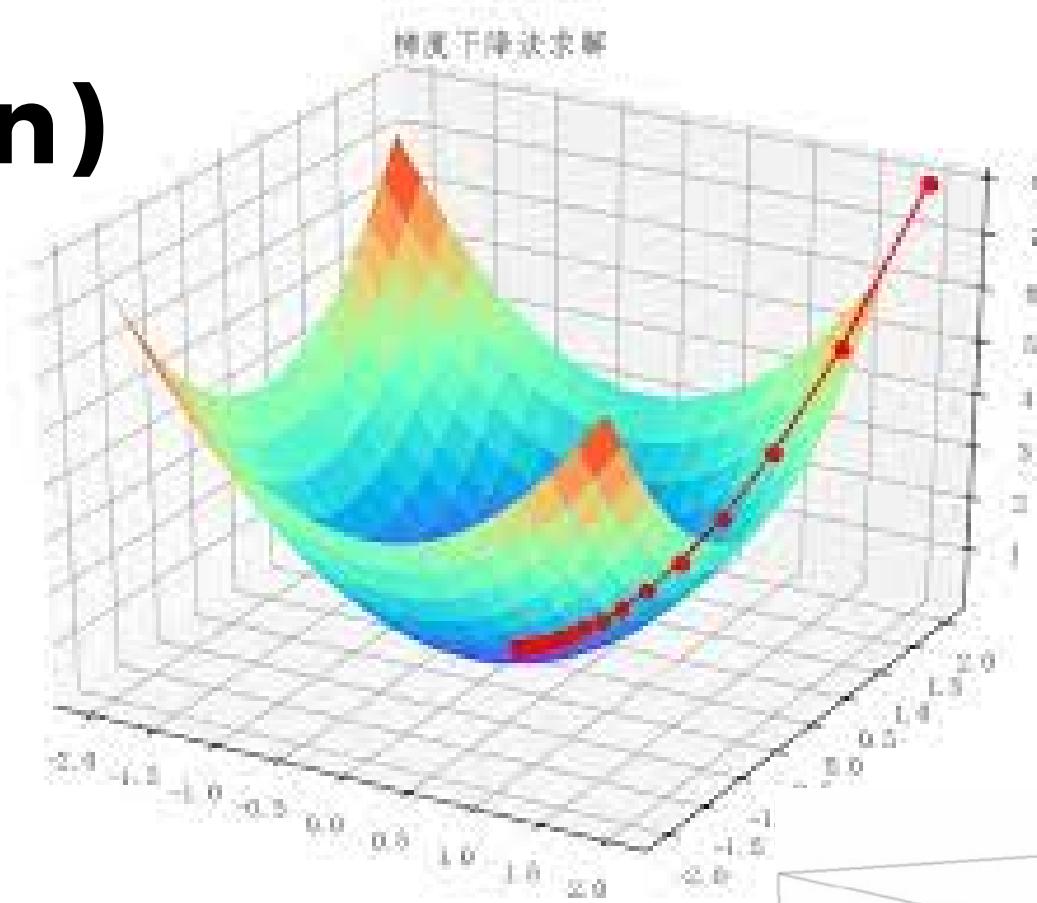
# Training (loss optimization)

$$\mathbf{W}^* = \operatorname{argmin}_{\mathbf{W}} \frac{1}{n} \sum_{i=1}^n \mathcal{L}(f(x^{(i)}; \mathbf{W}), y^{(i)})$$

$$\mathbf{W}^* = \operatorname{argmin}_{\mathbf{W}} J(\mathbf{W})$$

## Algorithm

1. Initialize weights randomly  $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Compute gradient,  $\frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
4. Update weights,  $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
5. Return weights

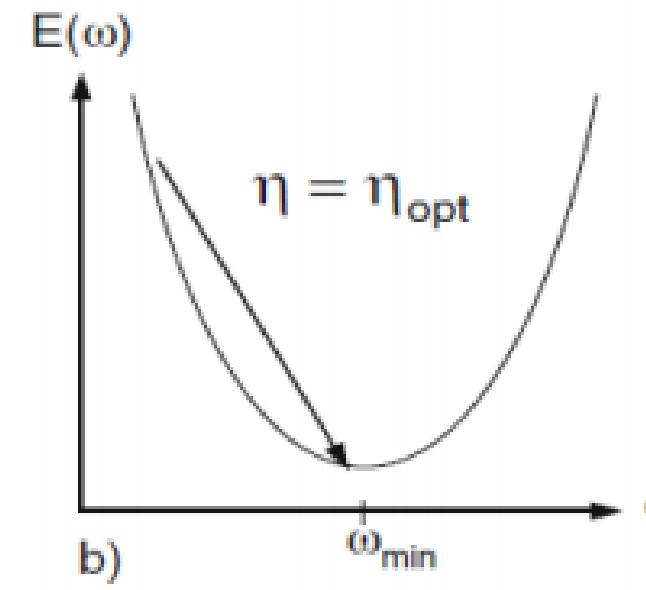
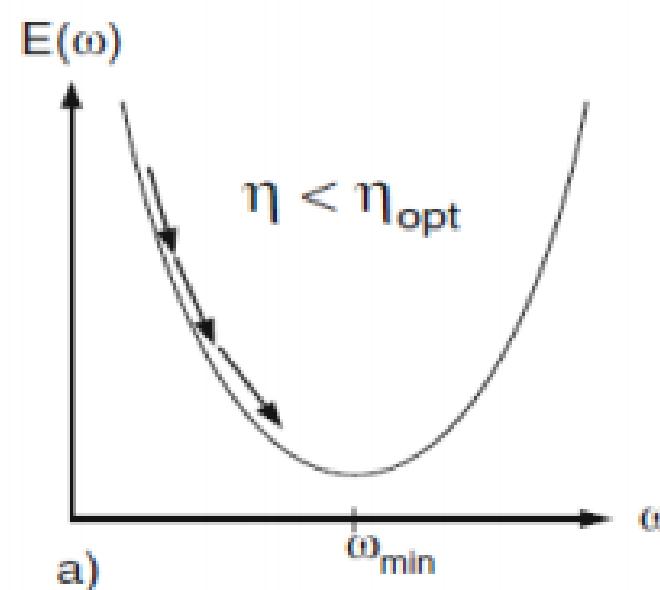


Principales problemas:

- Mínimos locales
- Silla de montar: el g. d. un vez que llega a una región con gradiente cero, no puede escapar de allí independientemente de la calidad del mínimo.

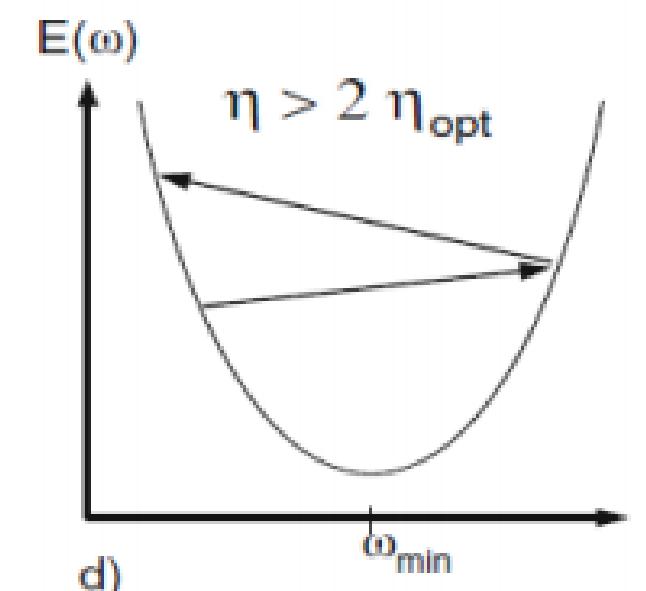
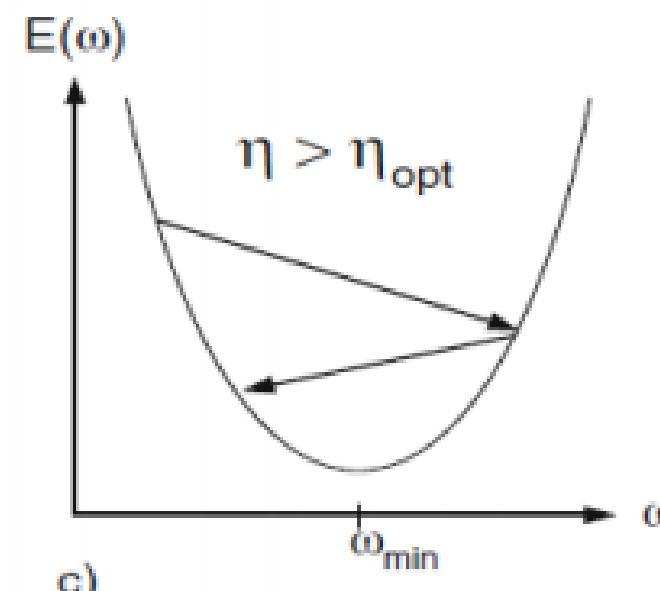


# Training (loss optimization)



$$\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$$

How can we set the learning rate?



- <https://ruder.io/optimizing-gradient-descent/index.html>

## Gradient Descent Algorithms

### Algorithm

- SGD
- Adam
- Adadelta
- Adagrad
- RMSProp

### TF Implementation

tf.keras.optimizers.SGD
tf.keras.optimizers.Adam
tf.keras.optimizers.Adadelta
tf.keras.optimizers.Adagrad
tf.keras.optimizers.RMSProp

### Reference

Kiefer & Wolfowitz. "Stochastic Estimation of the Maximum of a Regression Function." 1952.

Kingma et al. "Adam: A Method for Stochastic Optimization." 2014.

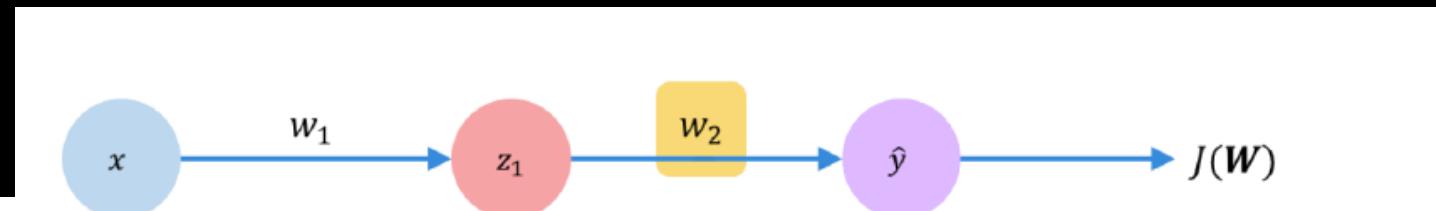
Zeiler et al. "ADADELTA: An Adaptive Learning Rate Method." 2012.

Duchi et al. "Adaptive Subgradient Methods for Online Learning and Stochastic Optimization." 2011.

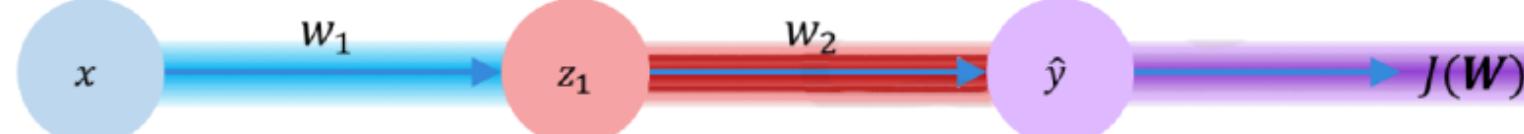


# Backpropagation

- Qué tanto puede afectar a la función de pérdida un pequeño cambio en uno de los pesos?



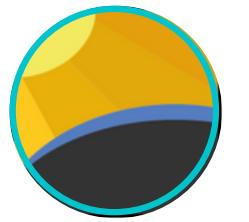
- Aplicando la regla de la cadena



$$\frac{\partial J(\mathbf{W})}{\partial w_1} = \underline{\frac{\partial J(\mathbf{W})}{\partial \hat{y}}} * \underline{\frac{\partial \hat{y}}{\partial z_1}} * \underline{\frac{\partial z_1}{\partial w_1}}$$

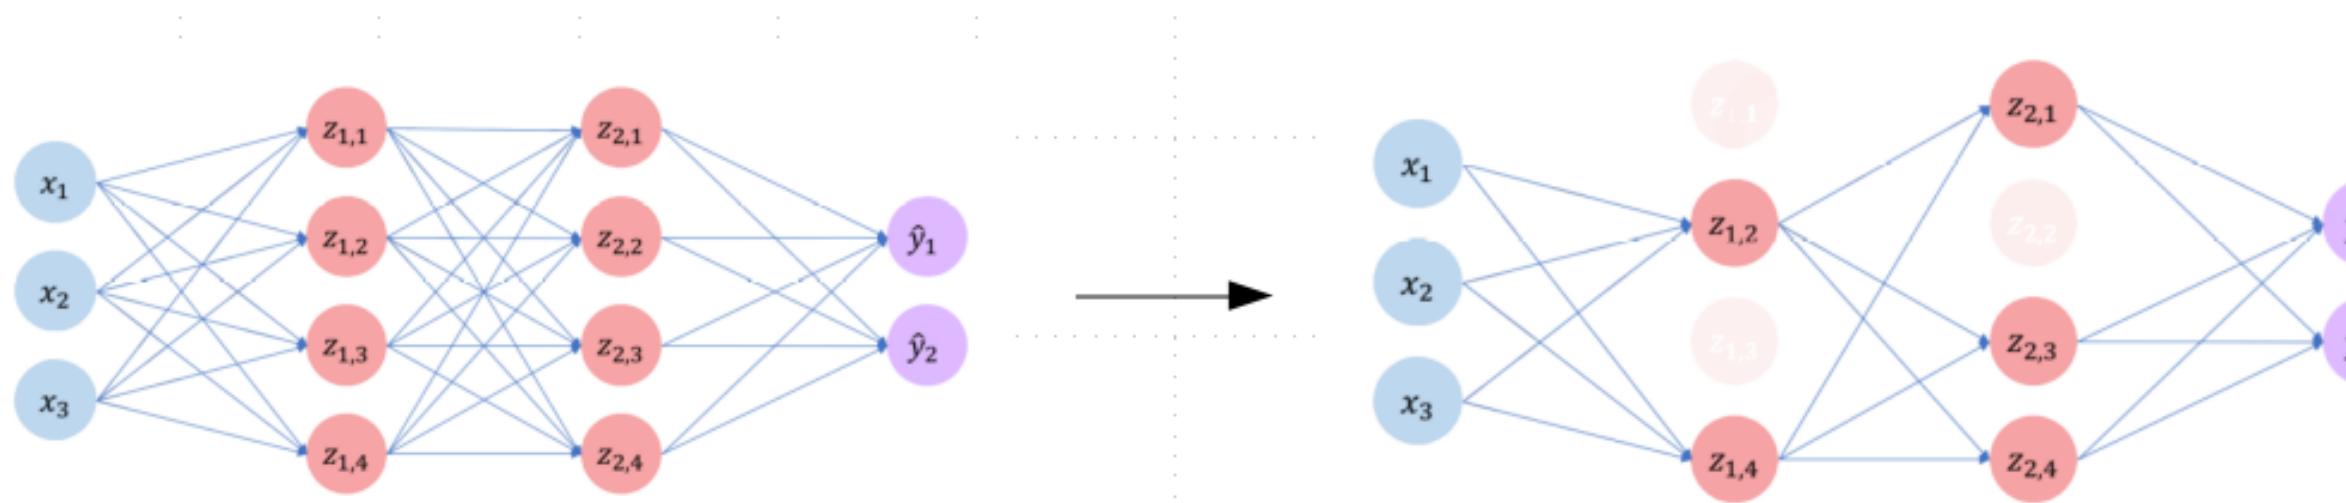
- Repetir para cada uno de los pesos en la red usando los gradientes de las capas siguientes!

- Info se propaga desde las entradas hacia adelante mediante sus parámetros hasta que logra hacer una predicción
- Luego realiza una propagación hacia atrás a lo largo de la red para ir modificando los parámetros de manera que el error final sea el mínimo.
- El error asociado a una mala predicción es usado para ajustar los parámetros (el aprendizaje puede ser visto como una optimización). Para esto la salida  $y$  es comparada con el valor esperado (target) del conjunto de entrenamiento ( $z$ ). La diferencia entre el valor esperado y el de salida se llama error o residuo.



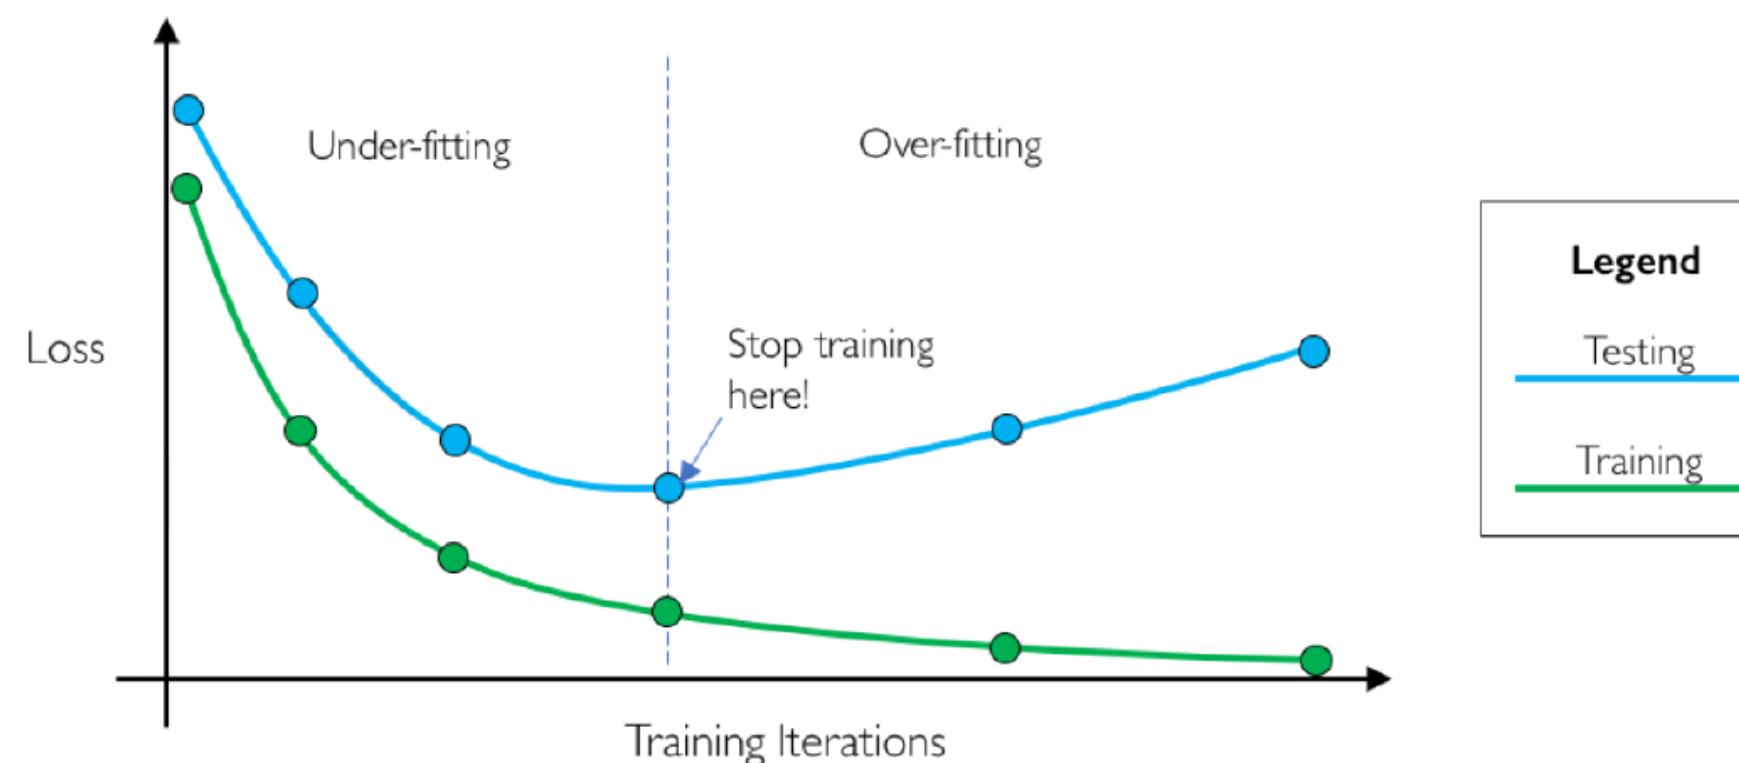
# Regularización

- Técnica que consiste en establecer restricciones al problema de optimización para evitar llegar a modelos complejos
- La idea es poder generalizar nuestro modelo para nuevos datos



## DROP OUT

- Durante el entrenamiento de manera aleatoria se setean algunas activaciones en 0.
- << overfitting

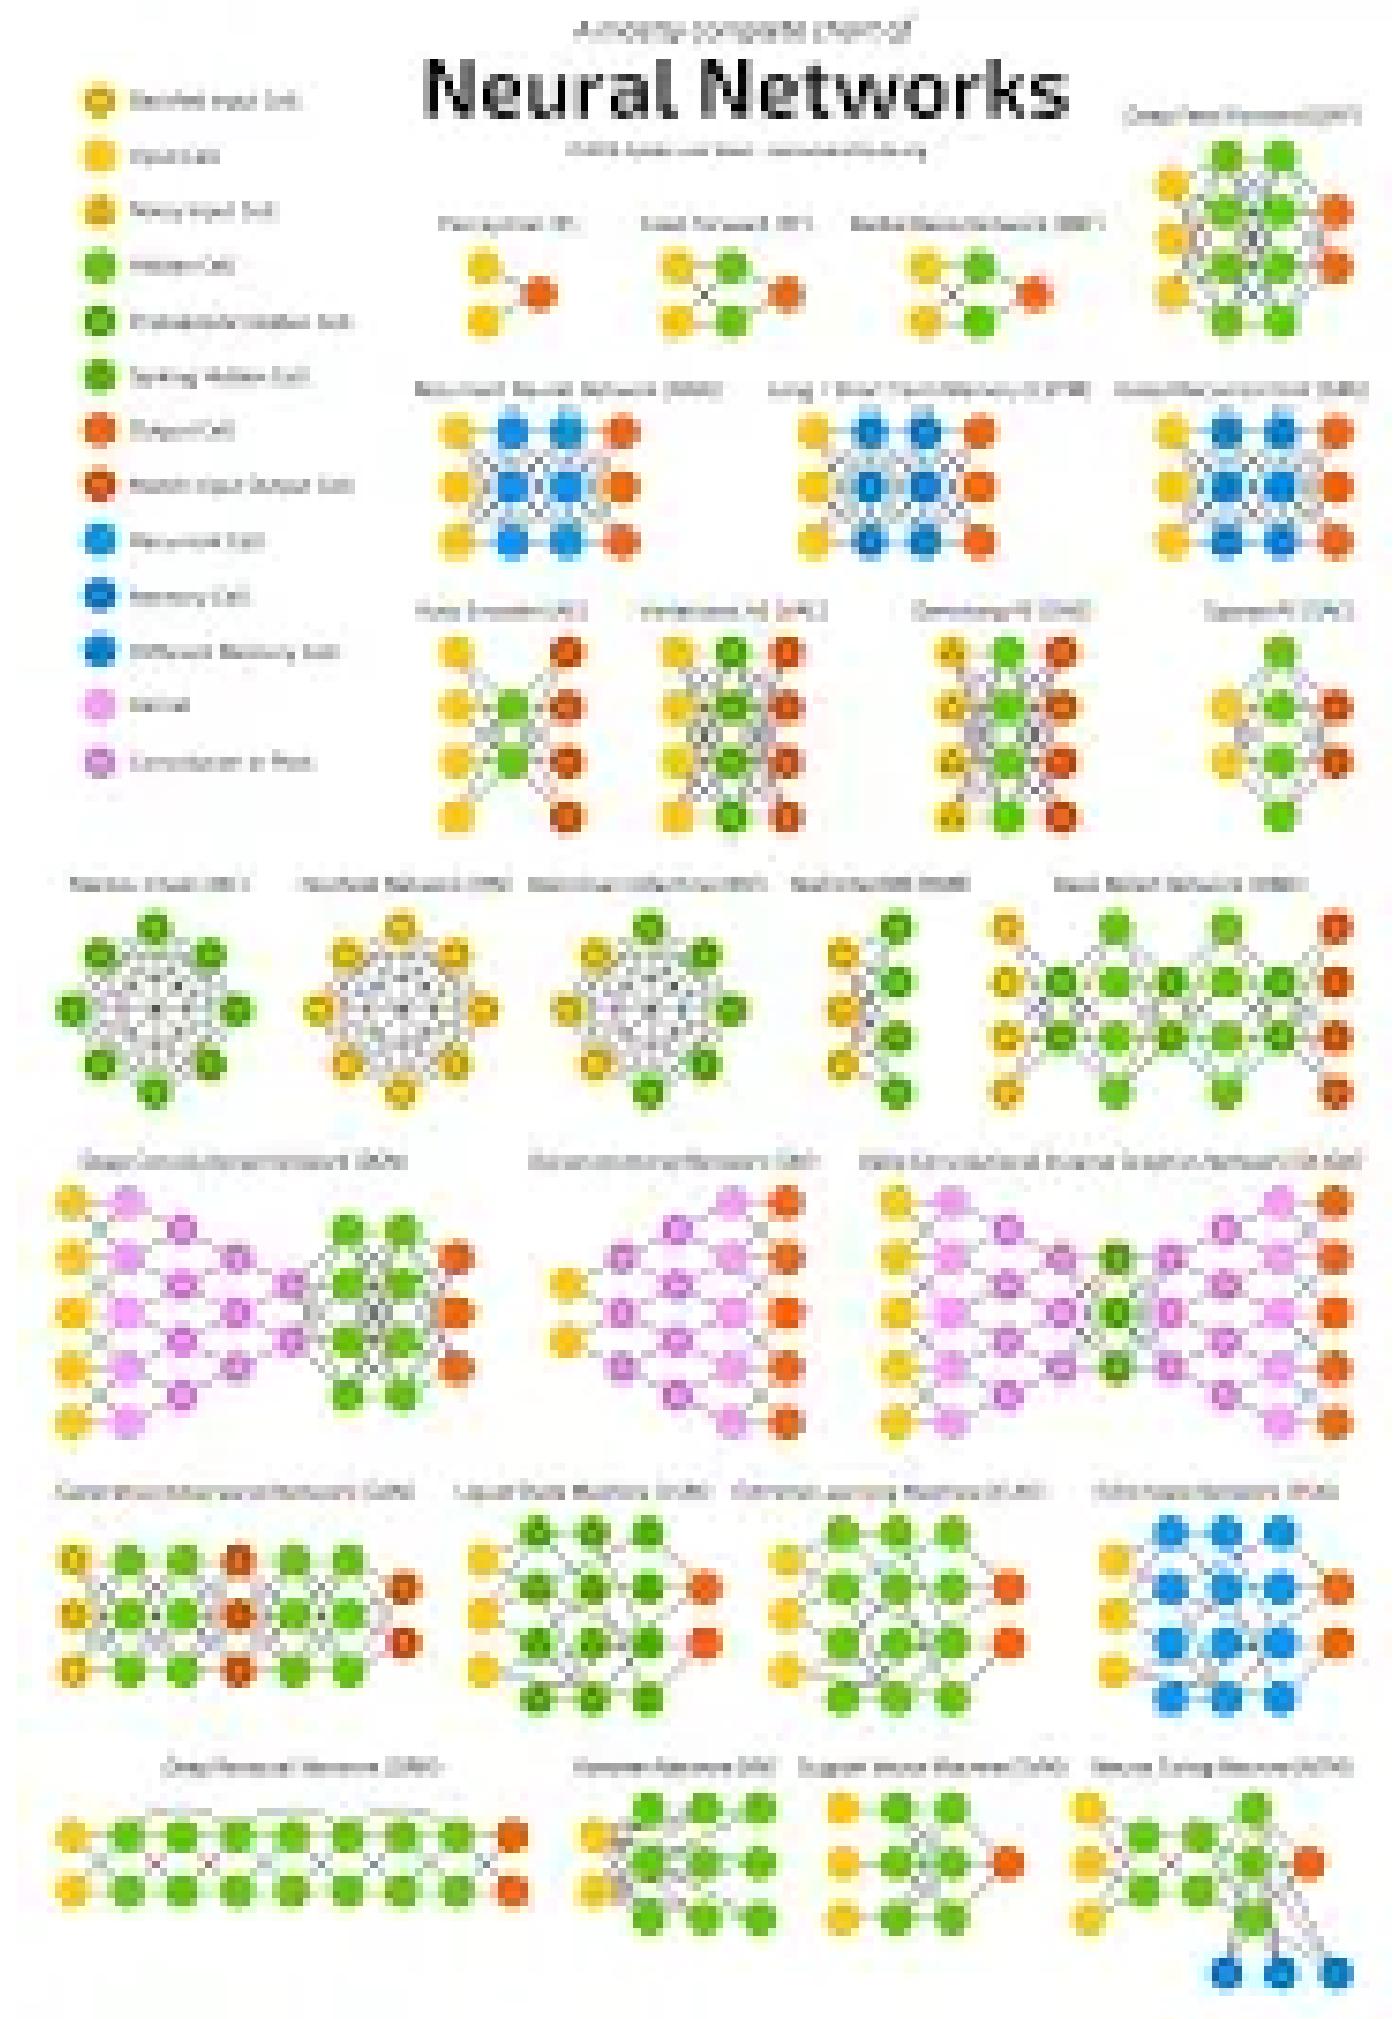


## EARLY STOPPING

- Consiste en parar antes de llegar al overfitting observando los datos de training y testing.



TSWC, 2022



```
32 self.file = None
33 self.fingerprints = {}
34 self.logdepth = 0
35 self.debug = False
36 self.logger = logging.getLogger('fingerprint')
37 if path:
38     self.file = open(path, 'w')
39 self.file.write('')
40 self.fingerprints = {}
41
42 @classmethod
43 def tree_settingscls(settings):
44     debug = settings.getbool('logger.debug')
45     return classmethod(_tree_settings(settings, debug))
46
47 def request_genuine(self, request):
48     fp = self.request_fingerprint(request)
49     if fp in self.fingerprints:
50         return True
51     self.fingerprints[fp] = []
52     self.file.write(fp + '\n')
53
54 def request_fingerprint(self, request):
55     return self.request_genuine(request)
```

# HANDS-ON LAB

## TP - 1

2 problems

- Radar signal classification (mandatory)
- Flare classification (optional)



TSWC, 2022



# From science to data science

Classify radar returns from the ionosphere as either suitable for further analysis or not (“good” or “bad”).

**PROBLEM 1 -** The 17 pairs of numbers, representing 17 discrete values of the real part and the corresponding 17 values of the complex part of an ACF, are the input to the neural network.

VINCENT G. SIGILLITO, SIMON P. WING, LARRIE V. HUTTON, and KILE B. BAKER

## CLASSIFICATION OF RADAR RETURNS FROM THE IONOSPHERE USING NEURAL NETWORKS

In ionospheric research, we must classify radar returns from the ionosphere as either suitable for further analysis or not. This time-consuming task has typically required human intervention. We tested several different feedforward neural networks to investigate the effects of network type (single-layer versus multilayer) and number of hidden nodes upon performance. As expected, the multilayer feedforward networks (MLFN's) outperformed the single-layer networks, achieving 100% accuracy on the training set and up to 98% accuracy on the testing set. Comparable figures for the single-layer networks were 94.5% and 92%, respectively. When measures of sensitivity, specificity, and proportion of variance accounted for by the model are considered, the superiority of the MLFN's over the single-layer networks is even more striking.

**Paper source:**

Sigillito, V. G., Wing, S. P., Hutton, L. V., \& Baker, K. B. (1989). Classification of radar returns from the ionosphere using neural networks. Johns Hopkins APL Technical Digest, 10, 262-266.

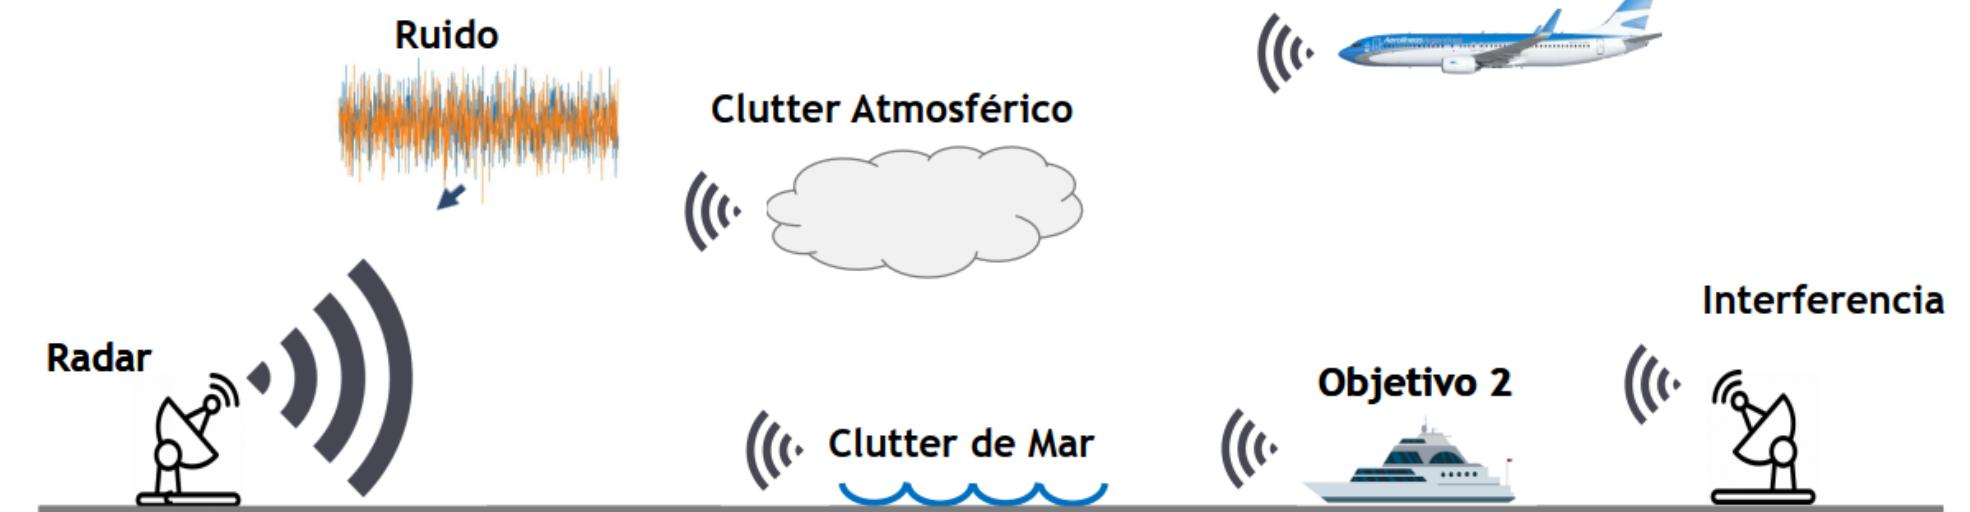
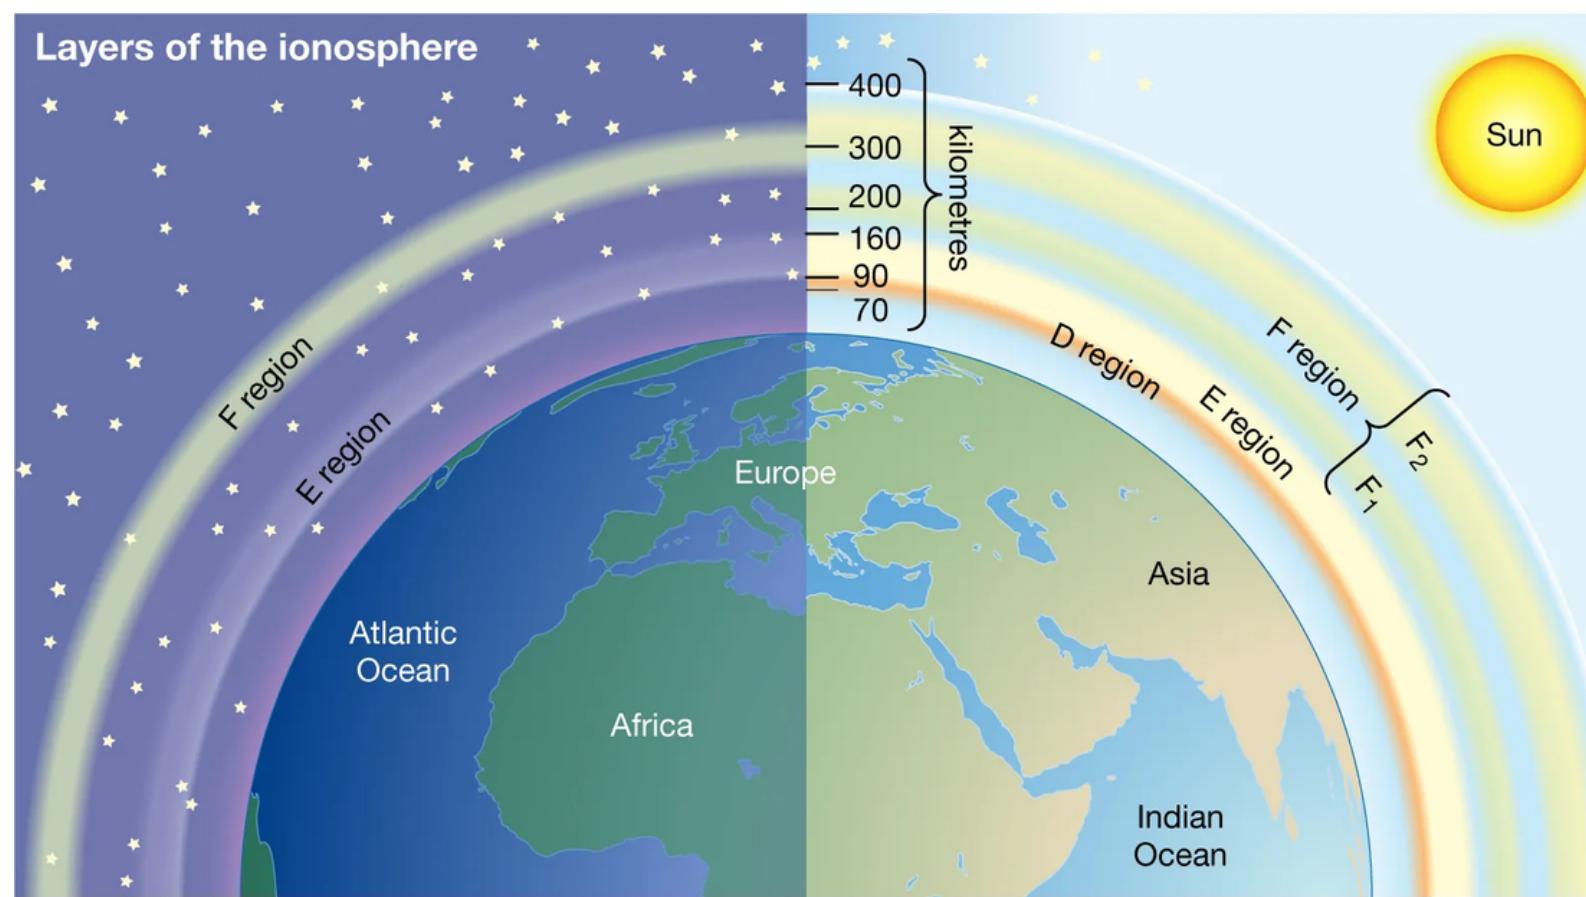
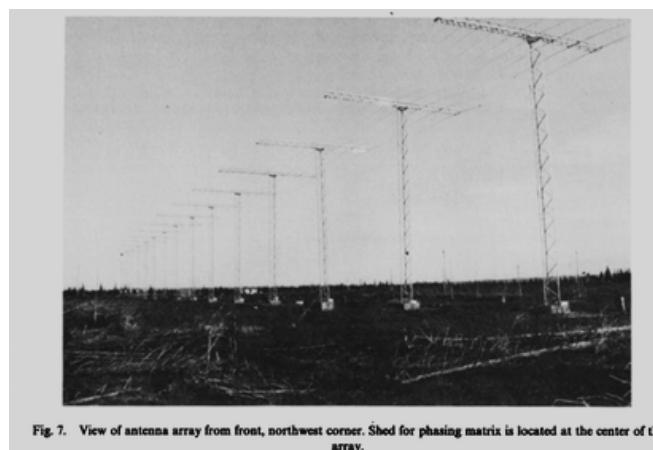
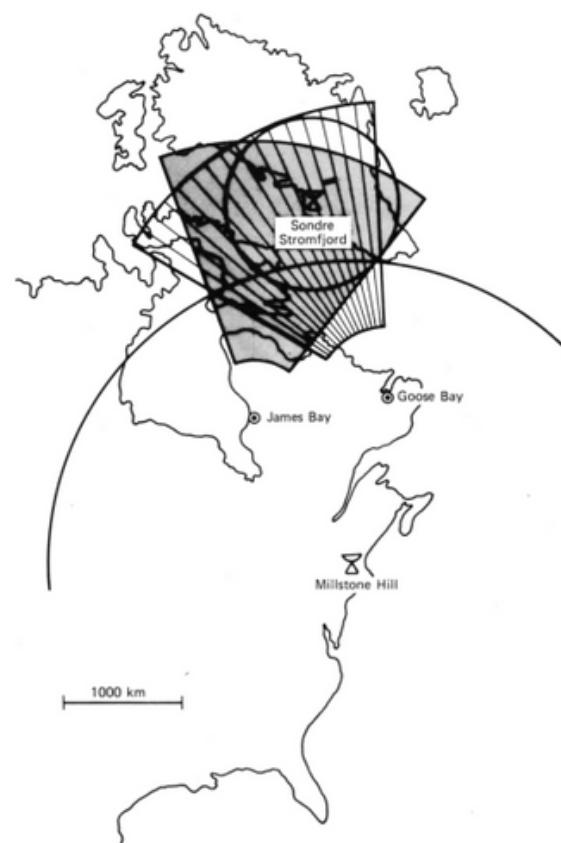
**Data source:**

Space Physics Group of the Johns Hopkins University Applied Physics Laboratory.



TSWC, 2022

# From science to data science



$$Eco(t) = Atte \cdot RCS \cdot u(t) \cdot m(t) \cdot e^{j\omega_c t \pm j\omega_D t + \varphi}$$

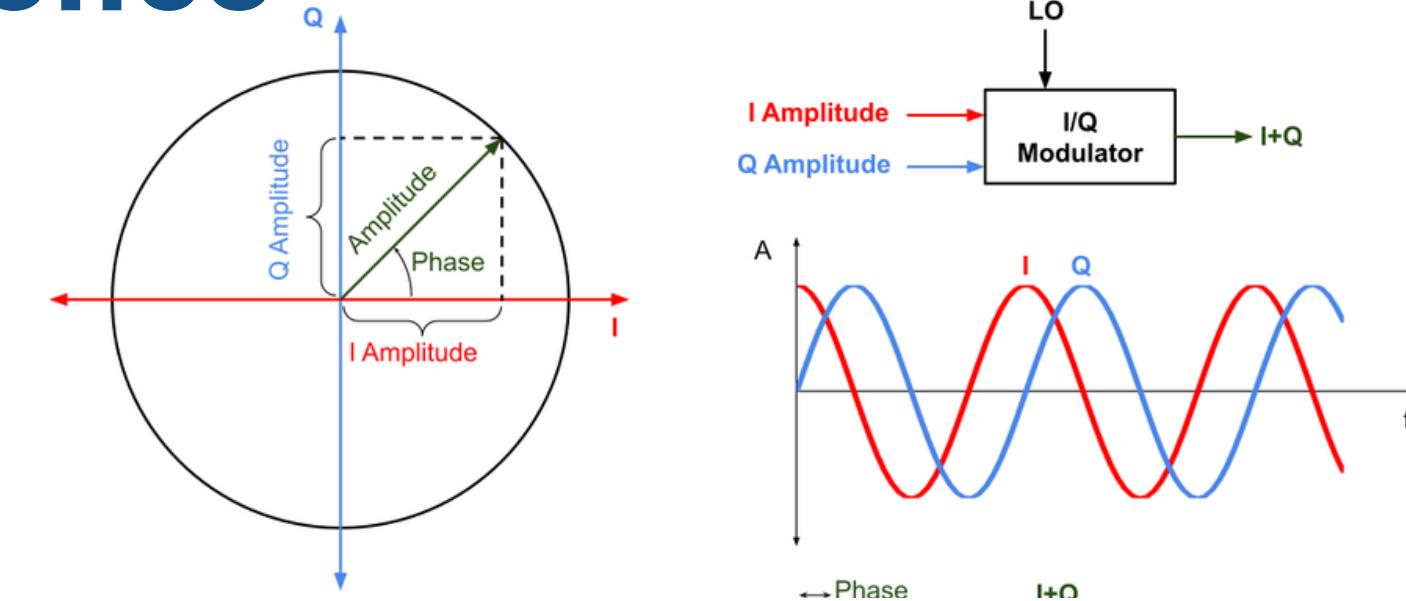
$$S_R(t) = Eco(t) + Clutter(t) + Ruido(t) + Interferencia(t)$$

\*: "target" = ionosphere

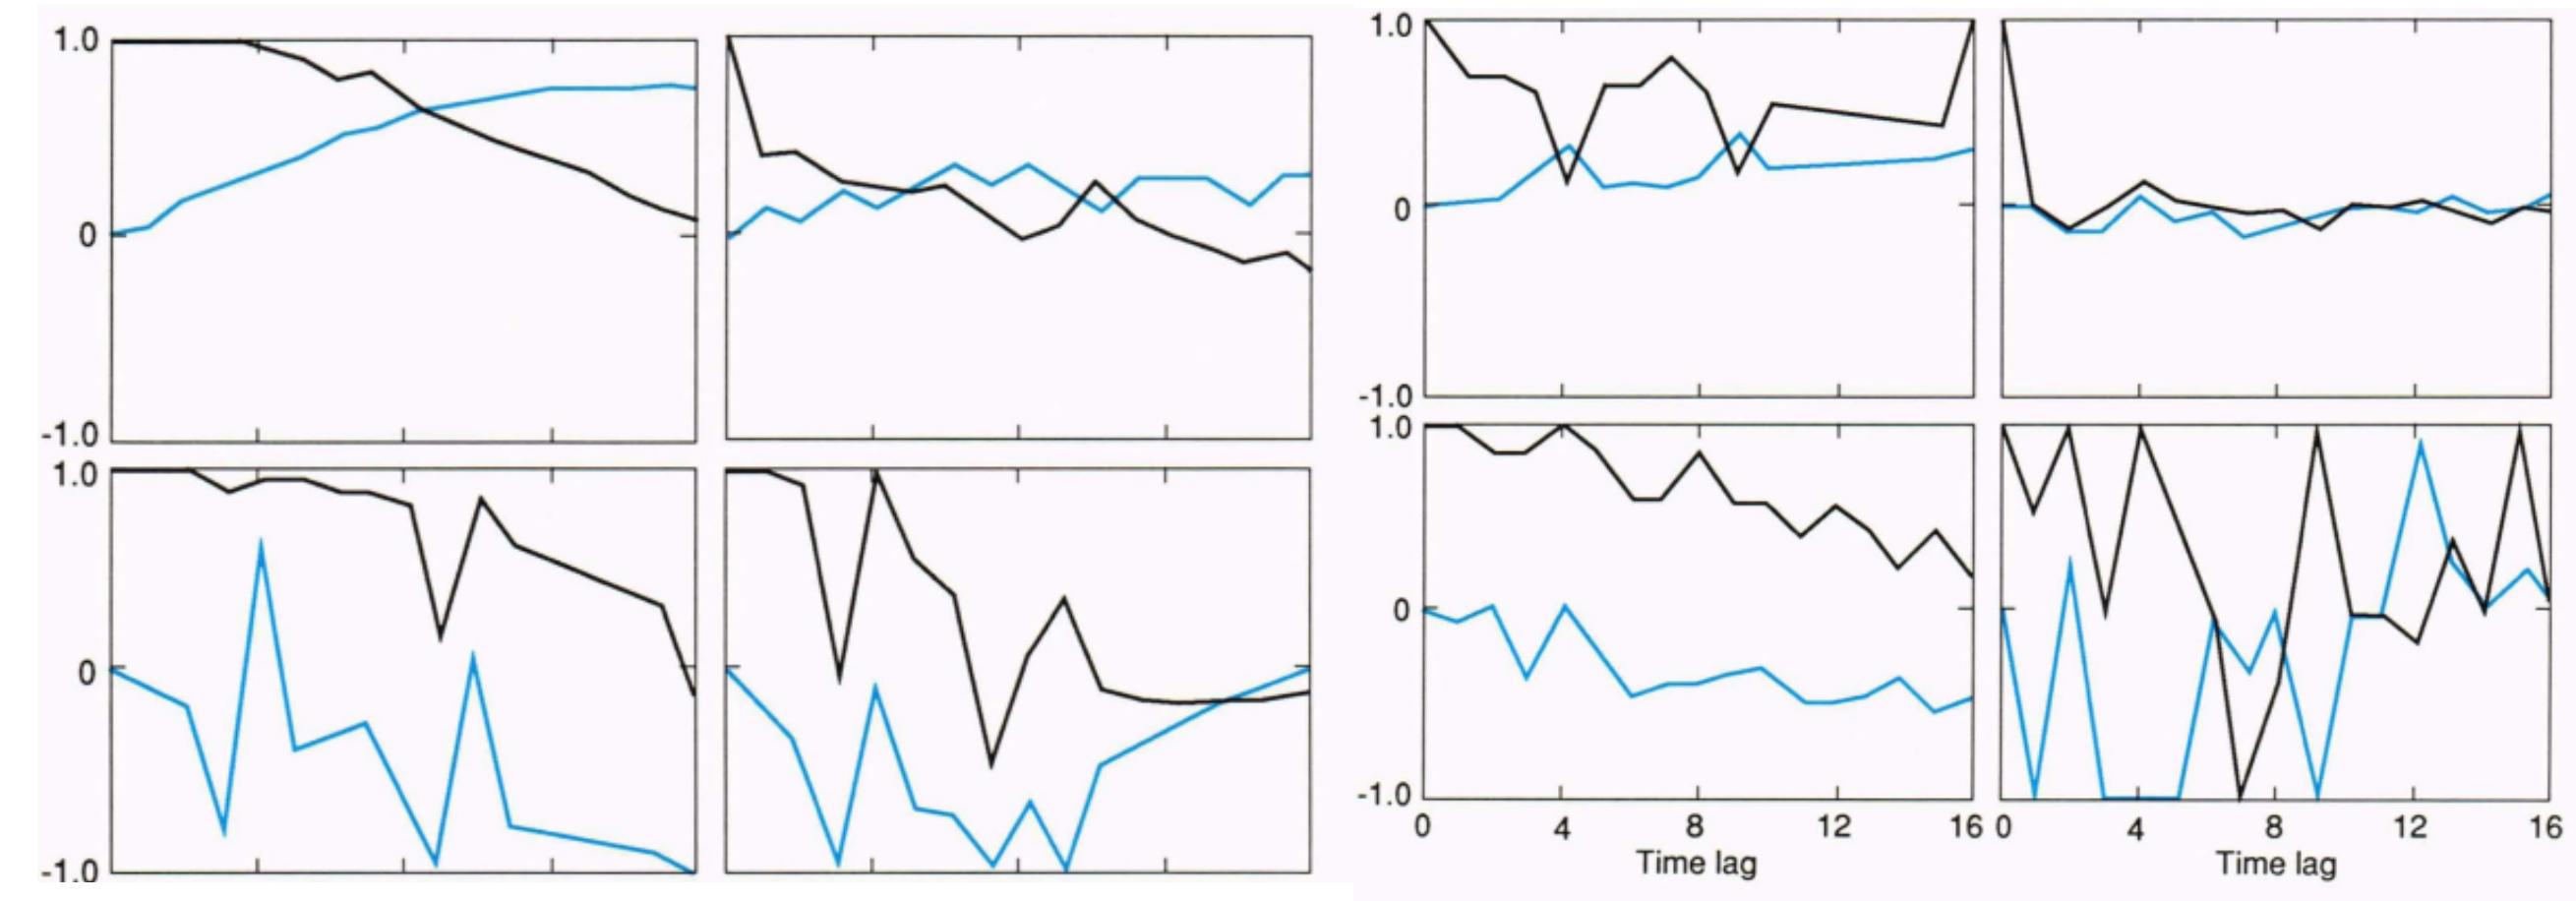


# From science to data science

**IQ Modulation:** it is a specific Phase-Modulation type, that provides us certain information in the real component of the signal and another possible coded information in the imaginary part. We can obtain both of them using demodulation techniques such as quadrature hybrids.



$$C(t) = A(t) + iB(t)$$



Good echo

Bad echo

Good echo

Bad echo

Auto-Correlation Function Output



# From science to data science



Copyright: SOHO/LASCO, SOHO/EIT (ESA & NASA)

28 Oct 2003 - X 17.2 flare -> SPE +  
CME (geoeffective, 29 Oct)

29 Oct 2003 - X 10.0 flare (same  
AR) -> CME

## The problem (optional)

- Automatic **classification of CMEs** according to their origin (flare/no-flare)
- From images (C2,C3 LASCO - SOHO) + expert observations
- source: [https://cdaw.gsfc.nasa.gov/CME\\_list/index.html](https://cdaw.gsfc.nasa.gov/CME_list/index.html)
- Acknowledgements: David Salazar de la Escuela de Física de la Universidad de Costa Rica + Salas-Matamoros & Klein, 2016.



# From science to data science

CME					Origen del CME									
					FLARE					Filamento				
Date		LASCO onset			Vprop	Coordinates		SXR onset			Observation date (Ha)		Coordinates	
DD	MM	YY	HH	MM		lat	long	HH	MM	DD	MM	YY	lat	long
7	4	97	14	27	592	-28	-11	14	8					
12	5	97	6	30	355	21	8	4	55					
21	6	98	5	35	619	18	39	5	18					
13	4	99	3	30	309	16	0	2	14					
6	6	0	15	54	1557	21	-10	15	26					
20	6	0	9	10	516	-27	38	8	27					
25	7	0	3	30	712	6	8	2	50					
21	7	6	13	54	403					20	7	6	-20	-25
29	12	6	21	30	236					29	12	6	-20	35
23	5	10	18	30	258					21	5	10	15	-20
30	1	11	19	36	521*					30	1	11	25	25
11	1	12	21	36	302					11	1	12	45	-25
16	3	12	20	36	862*					16	3	12	20	60
1	4	12	3	12	410*					31	3	12	20	-20
27	12	12	21	28	267					27	12	12	5	-20
.														
.														
.														

Columnas 1-3: Fecha de observación del CME

Columnas 4-5: Momento de la primera observación en C2 de LASCO reportado en el catálogo [https://cdaw.gsfc.nasa.gov/CME\\_list/](https://cdaw.gsfc.nasa.gov/CME_list/)

Columna 6: Velocidad de propagación del CME calculado mediante la fórmula empírica en Salas-Matamoros & Klein, 2016.

En eventos asociados a erupción de filamento, la velocidad de propagación del CME es la reportada en el catálogo [https://cdaw.gsfc.nasa.gov/CME\\_list/](https://cdaw.gsfc.nasa.gov/CME_list/), donde dicha velocidad es calculada por un ‘first-order polynomial fit’ a partir de mediciones de altura vrs tiempo.

Columnas 7-8: Localización de la región activa que dio origen al Flare.

Columnas 9-10: Momento de inicio del burst de SXR observado en los datos de GOES.

Columnas 11-13: Fechas de observación del filamento en imágenes en Ha (en [www.solarmonitor.org](http://www.solarmonitor.org)).

Columnas 14-15: La posición central de cada filamento.



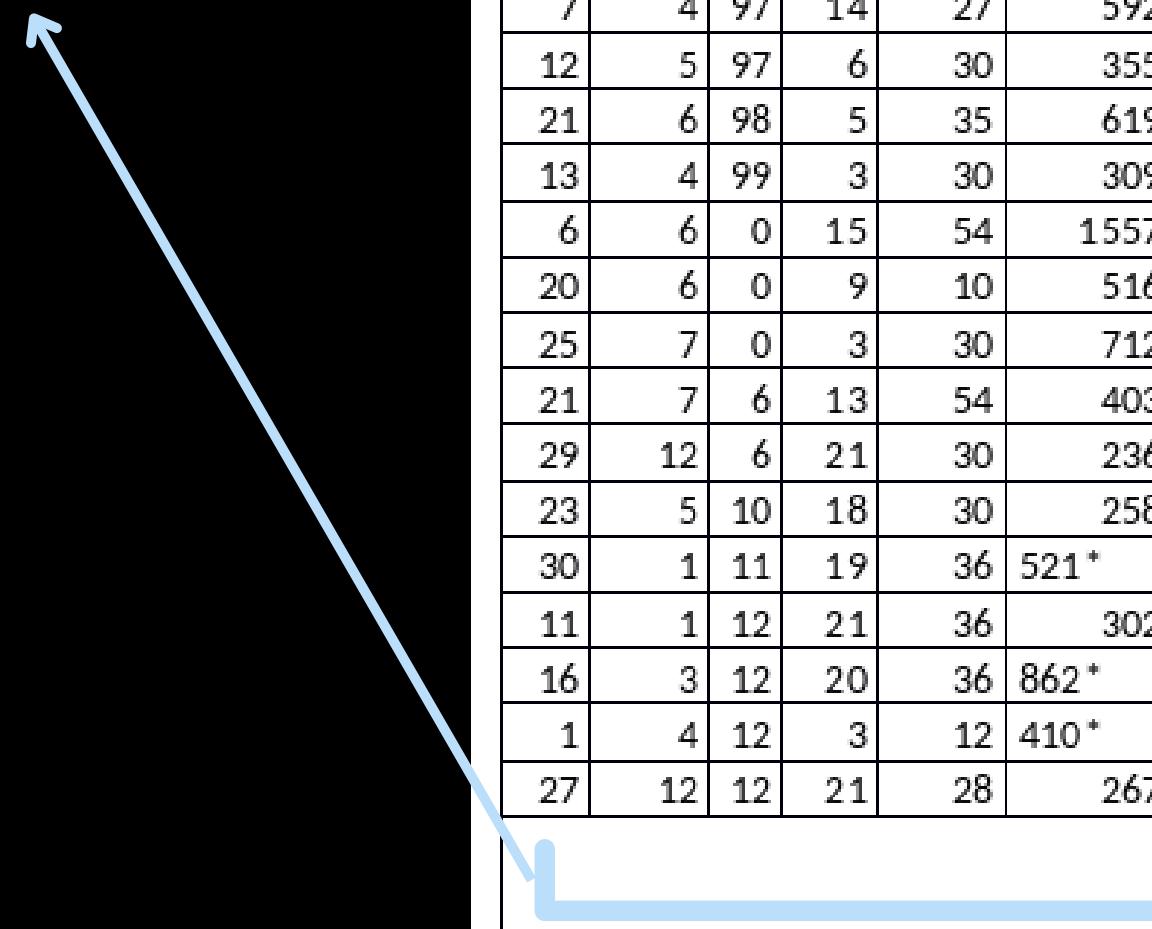
Data

Transformation



# From science to data science

- Features need selection & transformation



CME					Origen del CME									
Date			LASCO onset		Vprop	FLARE			Filamento			Observation date (Ha)		
DD	MM	YY	HH	MM		Coordinates	SXR onset	HH	MM	DD	MM	YY	lat	long
7	4	97	14	27	592	28	-11	14	8					
12	5	97	6	30	355	21	8	4	55					
21	6	98	5	35	619	18	39	5	18					
13	4	99	3	30	309	16	0	2	14					
6	6	0	15	54	1557	21	-10	15	26					
20	6	0	9	10	516	-27	38	8	27					
25	7	0	3	30	712	6	8	2	50					
21	7	6	13	54	403					20	7	6	-20	-25
29	12	6	21	30	236					29	12	6	-20	35
23	5	10	18	30	258					21	5	10	15	-20
30	1	11	19	36	521*					30	1	11	25	25
11	1	12	21	36	302					11	1	12	45	-25
16	3	12	20	36	862*					16	3	12	20	60
1	4	12	3	12	410*					31	3	12	20	-20
27	12	12	21	28	267					27	12	12	5	-20

Automatic **classification** of CMEs (supervised)

- Target needs transformation



# From science to data science

*Data Transformation*

CME					Vprop	Origen del CME											
Date		LASCO onset				Coordinates		SXR onset		Filamento			Observation date (Ha)		Coordinates		
DD	MM	YY	HH	MM		lat	long	HH	MM	DD	MM	YY	lat	long			
7	4	97	14	27		592	-28	-11	14	8							
12	5	97	6	30		355	21	8	4	55							
21	6	98	5	35		619	18	39	5	18							
13	4	99	3	30		309	16	0	2	14							
6	6	0	15	54		1557	21	-10	15	26							
20	6	0	9	10		516	-27	38	8	27							
25	7	0	3	30		712	6	8	2	50							
21	7	6	13	54		403				20	7	6	-20	-25			
29	12	6	21	30		236				29	12	6	-20	35			
23	5	10	18	30		258				21	5	10	15	-20			
30	1	11	19	36	521*					30	1	11	25	25			
11	1	12	21	36	302					11	1	12	45	-25			
16	3	12	20	36	862*					16	3	12	20	60			
1	4	12	3	12	410*					31	3	12	20	-20			
27	12	12	21	28	267					27	12	12	5	-20			



# From science to data science

- The training/test set

	A	B	C
1	SunActivity	vprop	Origin
2	low	592	flare
3	low	355	flare
4	low	619	flare
5	high	309	flare
6	high	1557	flare
7	high	516	flare
8	high	712	flare
9	high	214	flare

## To Do:

- data exploration ( outliers? nulls?, # samples, balanced? etc)
- data modeling preparation: standarization/normalization, one-hot encoding (or others), etc
- data splitting (70/30?, 80/20?, etc)
- MLP architecture! hyperparameters: #layers, Optimization algorithm, etc
- Fitting the model
- Obtain metrics
- Tunning

## Questions

- Which is the scores for the proposed model?
- How confident are we about the model performance/validation?
- Is the dataset representative enough
- propose other questions?!

# International Workshop on Machine Learning for Space Weather: Fundamentals, Tools and Future Prospects

**7-11 November 2022**  
**This is a hybrid meeting**  
**Buenos Aires, Argentina**



Further information:

<https://indico.ictp.it/event/9840/>  
smr3750@ictp.it  
+39-040-2240284  
Elizabeth Brancaccio

**Dra María Graciela Molina**  
FACET-UNT / CONICET  
Tucumán Space Weather Center - TSWC

<https://spaceweather.facet.unt.edu.ar/>  
IG -> @spaceweatherargentina

gmolina@herrera.unt.edu.ar

